

# Deep Learning

Power of Depth

Lionel Fillatre

2024-2025

# Outline

- One hidden-layer
- Limits of a single hidden layer
- Vanishing gradients
- Excess error
- Conclusion

One hidden-layer

# ReLU Deep Networks

- Let us remember that a feed-forward deep neural network has the form :

$$\hat{y} = f(x; \theta) = (f_{\theta_L} \circ f_{\theta_{L-1}} \circ \cdots \circ f_{\theta_1})(x)$$

where  $\theta = \{\theta_k : k = 1, \dots, L\}$  denotes the model parameters

- Let us assume that  $\theta_k = (W_k, b_k)$  and, for any  $h \in \mathbb{R}^{n_{k-1}}$ ,

$$f_{\theta_k}(h) = \sigma(W_k h + b_k)$$

where, for any  $z \in \mathbb{R}^n$ ,

$$\sigma(z) = \begin{pmatrix} \text{ReLU}(z_1) \\ \vdots \\ \text{ReLU}(z_n) \end{pmatrix} = \begin{pmatrix} \max\{0, z_1\} \\ \vdots \\ \max\{0, z_n\} \end{pmatrix}$$

# Approximation with ReLU nets


```
import numpy as np
import matplotlib.pyplot as plt

def relu(x):
    return np.maximum(x, 0)

def rect(x, a, b, h, eps=1e-7):
    return h / eps * (
        relu(x - a)
        - relu(x - (a + eps))
        - relu(x - b)
        + relu(x - (b + eps)))

x = np.linspace(-3, 3, 1000)
y = ( rect(x, -1, 0, 0.4))

plt.plot(x, y)
```



# Approximation with ReLU nets


```
import numpy as np
import matplotlib.pyplot as plt

def relu(x):
    return np.maximum(x, 0)

def rect(x, a, b, h, eps=1e-7):
    return h / eps * (
        relu(x - a)
        - relu(x - (a + eps))
        - relu(x - b)
        + relu(x - (b + eps)))

x = np.linspace(0.0, 5.0, 10) # 10 samples on the x-axis
z = np.linspace(0.0, 5.0, 500) # high resolution x-axis
sin_approx = np.zeros_like(z) # approximation of the sinusoid
for i in range(0, x.size-1):
    sin_approx = sin_approx + rect(z, x[i], x[i+1], 2+np.sin(x[i]))
plt.plot(z, 2+np.sin(z), 'r-')
plt.stairs(sin_approx[0:z.size-1], z)
plt.show()
```

Approximation of  $f(x) = 2 + \sin(x)$



# Approximation with ReLU nets


```
import numpy as np
import matplotlib.pyplot as plt

def relu(x):
    return np.maximum(x, 0)

def rect(x, a, b, h, eps=1e-7):
    return h / eps * (
        relu(x - a)
        - relu(x - (a + eps))
        - relu(x - b)
        + relu(x - (b + eps)))

x = np.linspace(0.0, 5.0, 200) # 200 samples on the x-axis
z = np.linspace(0.0, 5.0, 500) # high resolution x-axis
sin_approx = np.zeros_like(z) # approximation of the sinusoid
for i in range(0, x.size-1):
    sin_approx = sin_approx + rect(z, x[i], x[i+1], 2+np.sin(x[i]))
plt.plot(z, 2+np.sin(z), 'r-')
plt.stairs(sin_approx[0:z.size-1], z)
plt.show()
```

Approximation of  $f(x) = 2 + \sin(x)$



# Universal function approximation

**Theorem.** (Hornik et al, 1991)

- Let  $\sigma$  be a nonconstant, bounded, and monotonically-increasing continuous function.
- For any  $f \in C([0,1]^d)$  and  $\varepsilon > 0$ , there exists  $q \in \mathbb{N}$ , real constants  $v_i, b_i \in \mathbb{R}$  and  $w_i \in \mathbb{R}^d$  such that:

$$\left| \sum_{i=1}^q v_i \sigma(w_i^T x + b_i) - f(x) \right| < \varepsilon$$

- In other words, neural nets are dense in  $C([0,1]^d)$ .
- It guarantees that even a single hidden-layer network can represent any classification problem in which the boundary is locally linear (smooth);
- It does not inform about good/bad architectures, nor how they relate to the optimization procedure.
- The universal approximation theorem generalizes to any non-polynomial (possibly unbounded) activation function, including the ReLU (Leshno, 1993).



# Upper-Bound for one-hidden layer network

**Theorem** (Barron, 1994)

- The mean integrated square error  $\int \left( \hat{f}(x) - f(x) \right)^2 dx$  between the estimated network  $\hat{f}(x) = \sum_{i=1}^q v_i \sigma(w_i^T x + b_i) + v_0$  and the target function  $f$  is bounded by

$$O\left(\frac{C_f^2}{q}\right) + O\left(\frac{qn}{N} \log N\right)$$

where  $N$  is the number of training points,  $q$  is the number of neurons,  $n$  is the input dimension, and  $C_f$  measures the global smoothness of  $f$ .

- Provided enough data, it guarantees that adding more neurons will result in a better approximation.
- For your information,

$$C_f = \int \|\omega\|_1 \tilde{f}(\omega) d\omega$$


where  $\tilde{f}(\omega)$  is the Fourier transform of  $f(x)$ .

# Problem solved?

- Universal function approximation theorems do not tell us:
  - The number  $q$  of hidden units is small enough to have the network fit in RAM.
  - There is no constructive way to find an optimal solution
  - The optimal function parameters can be found in finite time by minimizing the Empirical Risk with SGD and the usual random initialization schemes.
- Going deeper?
  - Why would it be a good idea to stack more layers?
  - Are there any theoretical insights for doing this? Empirical ones?

# Limits of a single hidden layer

# Number of linear pieces



- Let  $\mathcal{F}$  be the set of piecewise linear mappings on  $[0,1]$
- Let  $\kappa(f)$  be the minimum number of linear pieces needed to represent  $f \in \mathcal{F}$ .
- Let  $\sigma: \mathbb{R} \rightarrow \mathbb{R}$  be the ReLU function

$$\sigma(x) = \text{ReLU}(x) = \max(0, x)$$

- If we compose  $\sigma$  and  $f \in \mathcal{F}$ , any linear piece that does not cross 0 remains a single piece or disappears, and one that does cross 0 breaks into two, hence

$$\forall f \in \mathcal{F}, \kappa(\sigma(f)) \leq 2 \kappa(f)$$

- We also have


$$\forall (f, g) \in \mathcal{F}^2, \kappa(f + g) \leq \kappa(f) + \kappa(g)$$

# Illustration



$$\kappa(f) = 16$$

ReLU( $x$ )



$$\kappa(\sigma(f)) = 14$$

# Bound on the number of linear pieces

- Consider a MLP with ReLU,  $D$  layers, a single input unit, and a single output unit.
  - Single unit input layer:  $h^0 = x \in \mathbb{R}^{n_0=1}$
  - Hidden layers:  $h^d = (h_1^d, \dots, h_{n_d}^d)$ ,  $\forall d = 1, \dots, D$ , with  $h_i^d = \sigma(z_i^{d-1})$  and  $z_i^{d-1} = \sum_{j=1}^{n_{d-1}} w_{ij}^{d-1} h_j^{d-1} + b_i^{d-1}$
  - Single unit output layer:  $\hat{y} = f(x) = h^D \in \mathbb{R}^{n_D=1}$

- Then, we get

$$\forall i, \ell, \quad \kappa(h_i^d) = \kappa(\sigma(z_i^{d-1})) \leq 2\kappa(z_i^{d-1}) \leq 2 \sum_{j=1}^{n_{d-1}} \kappa(h_j^{d-1})$$

- It follows that


$$\forall d, \quad \max_i \kappa(h_i^d) \leq 2n_{d-1} \max_j \kappa(h_j^{d-1})$$

- We get the following bound for any ReLU MLP

$$\kappa(\hat{y}) \leq 2^D \prod_{d=1}^D n_d$$


# Tight bound?

- Although this seems quite a pessimist bound, we can hand-design a network that (almost) reaches it



# Triangle wave


- So for any  $D$ , there is a network with  $D$  hidden layers and  $2D$  hidden units which computes a function  $f: [0,1] \rightarrow [0,1]$  of period  $\frac{1}{2^D}$





# Approximation of the triangle wave


- The unit segment  $[0,1]$  is cut into  $2^D$  segments



Is the proposed approximation (in blue) accurate?


# Approximation of the triangle wave

- Given  $g \in \mathcal{F}$  (blue curve), it crosses  $\frac{1}{2}$  at most  $\kappa(g)$  times,



# Approximation of the triangle wave


- On at least  $2^D - \kappa(g)$  segments of length  $\frac{1}{2^D}$ ,  $g$  is on one side of  $\frac{1}{2}$




# Approximation of the triangle wave

Area of the small grey right-angled triangle:

$$\frac{\frac{1}{2^D} \times \frac{1}{2}}{2} = \frac{1}{2^D} \times \frac{1}{16}$$



Area of the isosceles triangle:  $\frac{1}{2^D} \times \frac{1}{16} \times 2 = \frac{1}{2^D} \times \frac{1}{8}$



Area of the grey area:

$$1 \int_0^{\frac{1}{2^D}} \left| f(x) - \frac{1}{2} \right| dx = \frac{1}{2^D} \times \frac{1}{16} \times 4 = \frac{1}{2^D} \times \frac{1}{4}$$

# Approximation of the triangle wave $f(x)$

- Summary:

- Given  $g \in \mathcal{F}$ , it crosses  $\frac{1}{2}$  at most  $\kappa(g)$  times,
- Which means that, on at least  $2^D - \kappa(g)$  segments of length  $\frac{1}{2^D}$ , it is on one side of  $\frac{1}{2}$ ,

Error on 1 segment w.r.t. the constant function  $c(x) = \frac{1}{2}$ ,  
i.e., area of the isocles triangle in previous slides

- It follows that

$$\begin{aligned} \int_0^1 |f(x) - g(x)| dx &\geq (2^D - \kappa(g)) \times \overbrace{\frac{1}{2} \int_0^{\frac{1}{2^D}} \left| f(x) - \frac{1}{2} \right| dx}^{\text{Error on 1 segment}} \\ &= (2^D - \kappa(g)) \times \frac{1}{2^D} \times \frac{1}{8} \\ &= \frac{1}{8} \left( 1 - \frac{\kappa(g)}{2^D} \right) \end{aligned}$$

- We multiply  $f$  by 8 (and also  $g$ ) to get the final result:  $\int_0^1 |f(x) - g(x)| dx \geq 1 - \frac{\kappa(g)}{2^D}$

# ReLU MLPs with a single input/output

- There exists a network  $f$  with  $D$  layers, and  $2D$  internal units, such that, for any network  $g$  with  $D'$  layers of sizes  $\{n_1, \dots, n_{D'}\}$

$$\int_0^1 |f(x) - g(x)| dx \geq 1 - \frac{2^{D'}}{2^D} \prod_{d=1}^{D'} n_d$$


- In particular, with  $g$  a single hidden layer network ( $D' = 1$ )

$$\int_0^1 |f(x) - g(x)| dx \geq 1 - \frac{2n_1}{2^D}$$

- To approximate  $f$  properly, the width  $n_1$  of  $g$ 's hidden layer has to increase exponentially with  $f$ 's depth  $D$ .
- This is a simplified variant of results by Telgarsky (2015, 2016).

# Depth and Parametric Cost

- **Theorem** (Telgarsky, 2016): There exists functions that can be approximated by a deep ReLU network with  $\Theta(n^3)$  layers with  $\Theta(1)$  units that cannot be approximated by shallower networks with  $\Theta(n)$  layers unless they have  $\Omega(2^n)$  units.
- Note: the number of parameters of a deep network is typically quadratic with the number of units.
- This also holds for ReLU convnets with max pooling layers.
- Notation:
  - If a running time is  $\Omega(f(n))$ , then for large enough  $n$ , the running time is at least  $k \cdot f(n)$  for some constant  $C$
  - If a running time is  $\Theta(f(n))$ , then for large enough  $n$ , the running time is at most  $k \cdot f(n)$  for some constant  $C$



# The problem of depth

- Although it was known that deeper is better, for decades training deep neural networks was highly challenging and unstable.
- Besides limited hardware and data there were a few algorithmic flaws that have been fixed/softened in the last decade.
- An important issue is to control the amplitude of the gradient, which is tightly related to controlling activations.
- In particular we have to ensure that
  - The gradient does not « vanish » (Bengio et al., 1994; Hochreiter et al., 2001),
  - The gradient amplitude is homogeneous so that all parts of the network train at the same rate (Glorot and Bengio, 2010),
  - The gradient does not vary too unpredictably when the weights change (Balduzzi et al., 2017).




# What to do?

- Modern techniques change the functional itself instead of trying to improve training « from the outside » through penalty terms or better optimizers.
- Our main concern is to make the gradient descent work, even at the cost of engineering substantially the class of functions.
- An additional issue for training very large architectures is the computational cost, which often turns out to be the main practical problem.

# Vanishing gradients

# Vanishing gradients

- Training deep MLPs with many layers has for long (pre-2011) been very difficult due to the vanishing gradient problem.
  - Small gradients slow down, and eventually block, stochastic gradient descent.
  - This results in a limited capacity of learning.



Backpropagated gradients normalized histograms (Glorot and Bengio, 2010).  
Gradients for layers far from the output vanish to zero.

*Glorot and Bengio, Understanding the difficulty of training deep feedforward neural networks; AISTAT 2010*

# Vanishing gradients

- Consider a simplified 3-layer MLP, with  $x, w_1, w_2, w_3 \in \mathbb{R}$ , such that

$$f(x, w_1, w_2, w_3) = \sigma \left( w_3 \sigma \left( w_2 \sigma (w_1 x) \right) \right)$$

- Under the hood, this would be evaluated as

$$u_1 = w_1 x$$

$$u_2 = \sigma(u_1)$$

$$u_3 = w_2 u_2$$

$$u_4 = \sigma(u_3)$$

$$u_5 = w_3 u_4$$



$$\hat{y} = \sigma(u_5)$$

- Its derivative  $\frac{d\hat{y}}{dw_1}$  is

$$\frac{d\hat{y}}{dw_1} = \frac{d\hat{y}}{du_5} \frac{du_5}{du_4} \frac{du_4}{du_3} \frac{du_3}{du_2} \frac{du_2}{du_1} \frac{du_1}{dw_1} = \frac{\partial \sigma(u_5)}{\partial u_5} w_3 \frac{\partial \sigma(u_3)}{\partial u_3} w_2 \frac{\partial \sigma(u_1)}{\partial u_1} x$$

# Derivative of the sigmoid

- $\sigma(x) = \frac{1}{1+e^{-x}}$
- $\frac{d\sigma(x)}{dx} = \sigma(x)(1 - \sigma(x))$
- Hence, we get  
 $0 \leq \frac{d\sigma(x)}{dx} \leq \frac{1}{4}$  for all  $x$ .



# Bound on the derivative

- Assume that weights  $w_1, w_2, w_3$  are initialized randomly from a Gaussian with zero-mean and small variance, such that **with high probability  $-1 \leq w_i \leq 1$** .
- Then,

$$\left| \frac{d\hat{y}}{dw_1} \right| = \left| \frac{\partial \sigma(u_5)}{\partial u_5} \right| |w_3| \left| \frac{\partial \sigma(u_3)}{\partial u_3} \right| |w_2| \left| \frac{\partial \sigma(u_1)}{\partial u_1} \right| |x| \leq \left( \frac{1}{4} \right)^3 |x|$$

- This implies that the gradient  $\frac{d\hat{y}}{dw_1}$  **exponentially shrinks to zero** as the number of layers in the network increases. This is the vanishing gradient problem.
- In general, bounded activation functions (sigmoid, tanh, etc) are prone to the vanishing gradient problem.
- Note also the importance of a proper initialization scheme.



# Derivative of the ReLU function

- Note that the derivative of the ReLU function is

$$\frac{d\text{ReLU}(x)}{dx} = \begin{cases} 0 & \text{if } x \leq 0 \\ 1 & \text{if } x > 0 \end{cases}$$

- For  $x = 0$ , the derivative is undefined. In practice, it is set to zero.

$$\text{ReLU}(x) = \max(0, x)$$



# Solving the gradient vanishing problem?

- Assume again that weights  $w_1, w_2, w_3$  are initialized randomly from a Gaussian with zero-mean and small variance, such that with high probability  $-1 \leq w_i \leq 1$ .

- We get

$$\left| \frac{d\hat{y}}{dw_1} \right| = \left| \frac{\partial \sigma(u_5)}{\partial u_5} \right| |w_3| \left| \frac{\partial \sigma(u_3)}{\partial u_3} \right| |w_2| \left| \frac{\partial \sigma(u_1)}{\partial u_1} \right| |x| \leq |x|$$

- This solves the vanishing gradient problem, even for deep networks! (provided proper initialization)
- Note that:
  - The ReLU unit dies when its input is negative, which might block gradient descent.
  - This is actually a useful property to induce sparsity.
  - This issue can also be solved using leaky ReLUs, defined as

$$\text{LeakyReLU}(x) = \max(\alpha x, x)$$

for small  $\alpha > 0$



Excess error

# Learning (reminder)

- We observe some samples  $\mathcal{D} = (x_i, y_i)_{i=1, \dots, N}$  following the distribution of  $(X, Y)$ 
  - We are assuming that the samples are independent (iid assumption)
- We are considering a family  $\mathcal{F}$  of functions  $f_\theta$  parameterized by  $\theta \in \Theta$ 
  - The parameter  $\theta$  generally belongs to an Euclidean space (e.g.,  $\theta \in \mathbb{R}^n$ )
- We are expecting to solve

$$\theta^* \in \operatorname{argmin}_{\theta \in \Theta} \mathbb{E}(L(f_\theta(X), Y))$$

but, in practice, a naïve approach consists in minimizing the empirical risk

$$\hat{\theta} = \hat{\theta}(\mathcal{D}) \in \operatorname{argmin}_{\theta \in \Theta} \frac{1}{N} \sum_{i=1}^N L(f_\theta(x_i), y_i)$$

# Regularized Learning (recall)

- Generally, we prefer to solve

$$\hat{\theta} = \hat{\theta}(\mathcal{D}) \in \operatorname{argmin}_{\theta \in \Theta} \frac{1}{N} \sum_{i=1}^N L(f_{\theta}(x_i), y_i) + \lambda \Omega(f_{\theta})$$

where  $\Omega(f_{\theta})$  is the regularization term and  $\lambda \geq 0$  is an hyperparameter

- Examples:
  - Squared  $\ell_2$ -norm:  $\Omega(f_{\theta}) = \Omega(\theta) = \|\theta\|_2^2$
  - Sparsity:  $\Omega(f_{\theta}) = \Omega(\theta) = \|\theta\|_1$

# Decomposition of the error

- Let  $\mathcal{R}^* = \min_f \mathbb{E}(L(f(X), Y))$  the Bayes risk (minimum error)
- Let  $\mathcal{R}(\theta) = \mathbb{E}(L(f_\theta(X), Y))$  the expected risk at parameter  $\theta$  for function  $f_\theta \in \mathcal{F}$
- Let  $\hat{\mathcal{R}}(\theta) = \frac{1}{N} \sum_{i=1}^N L(f_\theta(x_i), y_i)$  the empirical risk at parameter  $\theta$
- Let  $\hat{\theta} = \hat{\theta}(\mathcal{D}) \in \operatorname{argmin}_{\theta} \hat{\mathcal{R}}(\theta)$
- We get the famous equality

$$\mathcal{R}(\hat{\theta}) - \mathcal{R}^* = \left( \mathcal{R}(\hat{\theta}) - \min_{\theta \in \Theta} \mathcal{R}(\theta) \right) + \left( \min_{\theta \in \Theta} \mathcal{R}(\theta) - \mathcal{R}^* \right)$$



Excess error

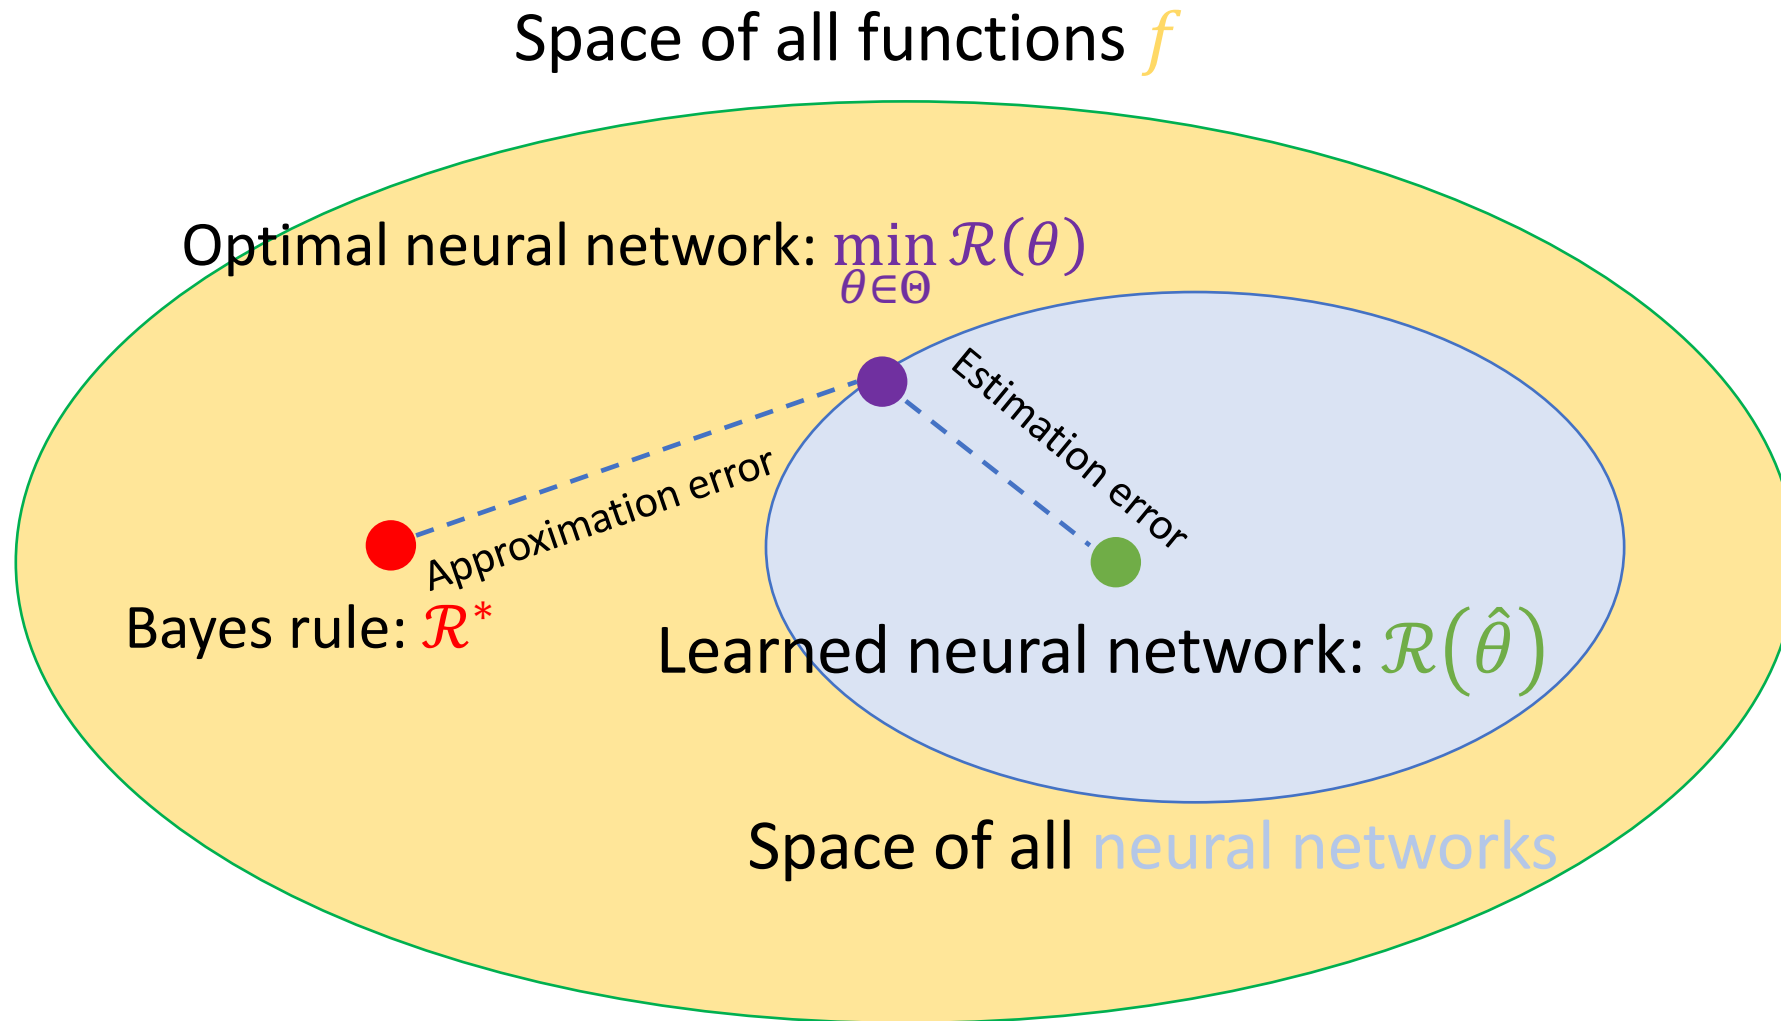


Estimation error



Approximation error

# Decomposition of the error



# In practice, decomposition of the error

- Let us assume that our minimization algorithm returns an approximate solution  $\tilde{\theta}$  such that

$$\hat{\mathcal{R}}(\tilde{\theta}) \leq \hat{\mathcal{R}}(\hat{\theta}) + \varepsilon$$

where  $\varepsilon \geq 0$  is a predefined tolerance (early stopping for example)

- Then, we get the equality

$$\varepsilon = \underbrace{\mathcal{R}(\tilde{\theta}) - \mathcal{R}^*}_{\text{Excess error}} = \underbrace{\left(\mathcal{R}(\hat{\theta}) - \min_{\theta \in \Theta} \mathcal{R}(\theta)\right)}_{\text{Estimation error}} + \underbrace{\left(\min_{\theta \in \Theta} \mathcal{R}(\theta) - \mathcal{R}^*\right)}_{\text{Approximation error}} + \underbrace{\left(\mathcal{R}(\tilde{\theta}) - \mathcal{R}(\hat{\theta})\right)}_{\text{Optimization Error}}$$

- The optimization is generally small with respect to the other errors
- Take home message: the computation of the parameters may have an impact of the classifier.

# Interpretation

- The approximation error measures how closely functions in  $\mathcal{F}$  can approximate the optimal solution  $f^b$ .
- The estimation error measures the effect of minimizing the empirical risk  $\hat{\mathcal{R}}(\theta)$  instead of the expected risk  $\mathcal{R}(\theta)$ .
  - The estimation error is determined by the number of training examples and by the capacity of the family of functions (see next slides).
- Large families of functions  $\mathcal{F}$  have smaller approximation errors but lead to higher estimation errors.
- This tradeoff has been extensively discussed and lead to excess error that scale between the inverse  $\left(\frac{1}{N}\right)$  and the inverse square root  $\left(\frac{1}{\sqrt{N}}\right)$  of the number of examples.

# Analysis of the Approximation error

- Approximation error:

$$\min_{\theta \in \Theta} \mathcal{R}(\theta) - \mathcal{R}^* = \min_{f_{\theta} \in \mathcal{F}} \mathbb{E}(L(f_{\theta}(X), Y)) - \min_f \mathbb{E}(L(f(X), Y))$$

- Main properties:
  - This error is non-random
  - It depends on the richness of  $\mathcal{F}$
  - If the optimal classifier belongs to  $\mathcal{F}$ , this error vanishes
  - This error is related to the functional analysis (approximation of function)



# Analysis of the Estimation error

- Warning:
  - $\hat{\theta}$  and  $\hat{\mathcal{R}}$  are random so all the terms depending on them (or just one of them) are random

- We have

$$\mathcal{R}(\hat{\theta}) - \mathcal{R}(\theta^*) = \mathcal{R}(\hat{\theta}) - \hat{\mathcal{R}}(\hat{\theta}) + \hat{\mathcal{R}}(\hat{\theta}) - \mathcal{R}(\theta^*)$$

- Note that  $\hat{\mathcal{R}}(\hat{\theta}) \leq \hat{\mathcal{R}}(\theta^*)$ . Hence,

$$\mathcal{R}(\hat{\theta}) - \mathcal{R}(\theta^*) \leq \mathcal{R}(\hat{\theta}) - \hat{\mathcal{R}}(\hat{\theta}) + \hat{\mathcal{R}}(\theta^*) - \mathcal{R}(\theta^*) \leq \mathcal{R}(\hat{\theta}) - \hat{\mathcal{R}}(\hat{\theta}) + \sup_{\theta \in \Theta} |\mathcal{R}(\theta) - \hat{\mathcal{R}}(\theta)|$$

- It follows that

$$\mathcal{R}(\hat{\theta}) - \mathcal{R}(\theta^*) \leq 2 \sup_{\theta \in \Theta} |\mathcal{R}(\theta) - \hat{\mathcal{R}}(\theta)|$$

# Generalization error

- The term  $|\mathcal{R}(\theta) - \hat{\mathcal{R}}(\theta)|$  is called the **generalization error**: it is the difference between the expected loss and the empirical loss
- Affine classifier:  $f_{\theta}(x) = \theta^T x = \sum_{i=1}^n \theta_i x_i$ 
  - $|\mathcal{R}(\theta) - \hat{\mathcal{R}}(\theta)| \approx \sqrt{\frac{n}{N}}$
- One-layer ReLU networks with regularization:  $f_{\theta}(x) = \sum_{i=1}^q v_i \max(0, w_i^T x + b_i)$ 
  - $|\mathcal{R}(\theta) - \hat{\mathcal{R}}(\theta)| \approx q \sqrt{\frac{n}{N}}$  where  $q$  is the number of hidden neurons


# Growth function, VC-dimension, shattering

- Let  $\mathcal{F}$  denote a class of functions from  $\mathcal{X}$  to  $\{0,1\}$  (the classification rules).
- For any non-negative integer  $m$ , we define the growth function of  $\mathcal{F}$  as

$$\Pi_{\mathcal{F}}(m) = \max_{x_1, \dots, x_m \in \mathcal{X}} |\{(f(x_1), \dots, f(x_m)) : f \in \mathcal{F}\}|$$

- Example: see the next slide!
- If  $|\{(f(x_1), \dots, f(x_m)) : f \in \mathcal{F}\}| = 2^m$ , we say  $\mathcal{F}$  shatters the set  $\{x_1, \dots, x_m\}$ .
  - Example: see the next slide!
- The Vapnik-Chervonenkis dimension of  $\mathcal{F}$  is denoted  $\text{VCdim}(\mathcal{F})$ 
  - It is the size of the largest shattered set, i.e. the largest  $m$  such that  $\Pi_{\mathcal{F}}(m) = 2^m$ .
  - If there is no largest  $m$ , we define  $\text{VCdim}(\mathcal{F}) = \infty$ .

# Example of growth function




- Set of step classifiers (threshold classifier):  $\mathcal{F} = \{f \in \{0,1\}^{\mathbb{R}} : f(x) = 1 \Leftrightarrow x > a, a \in \mathbb{R}\}$
- Classification results:
  - $\{(f(x_1), f(x_2), f(x_3)) : f \in \mathcal{F}\} = \{(0,0,0), (0,0,1), (0,1,1), (1,1,1)\}$
  - $|\{(f(x_1), f(x_2), f(x_3)) : f \in \mathcal{F}\}| = 4 < 2^3 = 8$
- $\mathcal{F}$  **does not shatter** the set  $\{x_1, x_2, x_3\}$ 
  - It is impossible to get the classification result  $(1,0,1)$  with a function  $f$  in  $\mathcal{F}$
- $\mathcal{F}$  **does not shatter** the set  $\{x_1, x_2\}$
- $\mathcal{F}$  **shatters** the set  $\{x_1\}$  whatever  $x_1$  is. Hence  $\text{VCdim}(\mathcal{F}) = 1$ .


# Example of VC dimension

- For an input space with two variables and a linear classifier, the VC dimension is 3.

A linear classifier can always shatter 3 points



A linear classifier may not shatter 4 points



# PAC Bound

- In case of binary classification, we have for example a more precise probabilistic behavior: the « Probably and Approximately Correct » (PAC) bound
- For any set  $\mathcal{F}$  with  $\text{VCdim}(\mathcal{F}) = d$ , with probability  $1 - \delta$  over the training dataset with  $N$  samples, we have

$$\mathcal{R}(\theta) \leq \hat{\mathcal{R}}(\theta) + O\left(\sqrt{\frac{d}{N} \log\left(\frac{N}{d}\right) - \frac{1}{N} \log(\delta)}\right)$$

# Finite-sample expressivity

- As soon as the number of parameters of a network is greater than  $N$ , even simple two-layer neural networks can represent any function of the input sample.
- We say that a neural network  $f_\theta(x)$  can represent any function of a sample of size  $N$  in  $n$  dimensions if for every sample  $S \subset (\mathbb{R}^n)^N$  with  $|S| = N$  and every function  $f: S \rightarrow \mathbb{R}$ , there exists a setting of the weights  $\theta$  of  $f_\theta(x)$  such that  $f_\theta(x) = f(x)$  for every  $x \in S$ .

## **Theorem** (Zhang, 2016)

- There exists a two-layer neural network with ReLU activations and  $2N + n$  weights that can represent any function on a sample of size  $N$  in  $n$  dimensions.

## **Corollary**

- For every  $k \geq 2$ , there exists a neural network with ReLU activations of depth  $k$ , width  $O\left(\frac{N}{k}\right)$  and  $O(N + n)$  weights that can represent any function on a sample of size  $N$  in  $n$  dimensions.

# Conclusion



# Conclusion

- Neural networks with several layers provide high capability for approximating multivariate functions
- Depth leads to vanishing gradient, which is an important issue
- Activation functions play an important role
- Theoretical results with neural network representations are in progress.