

# Deep Learning

Transformer

Lionel Fillatre

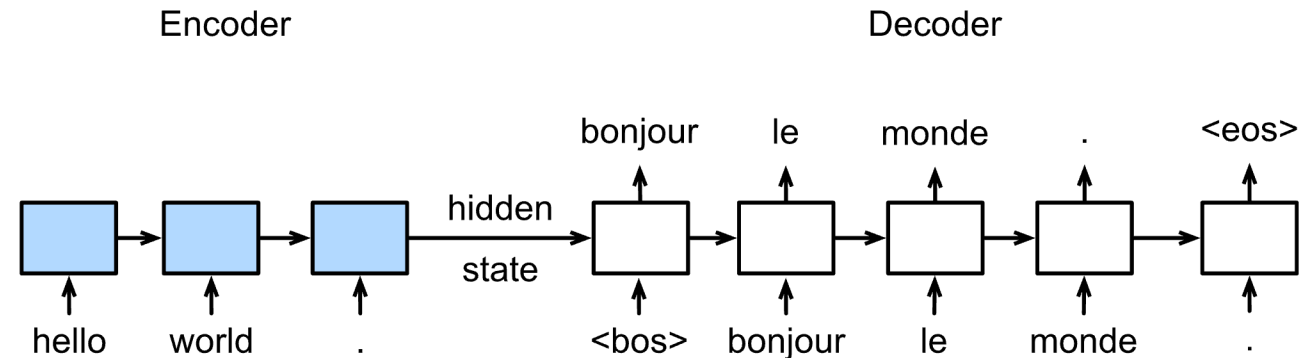
2024-2025

# Outline

- Sequence to Sequence
- Attention Mechanism
- Transformer's Encoder
- Input Encoding
- Classification Transformer
- Transformer for Sequence Transduction
- Conclusion

# Sequence to Sequence

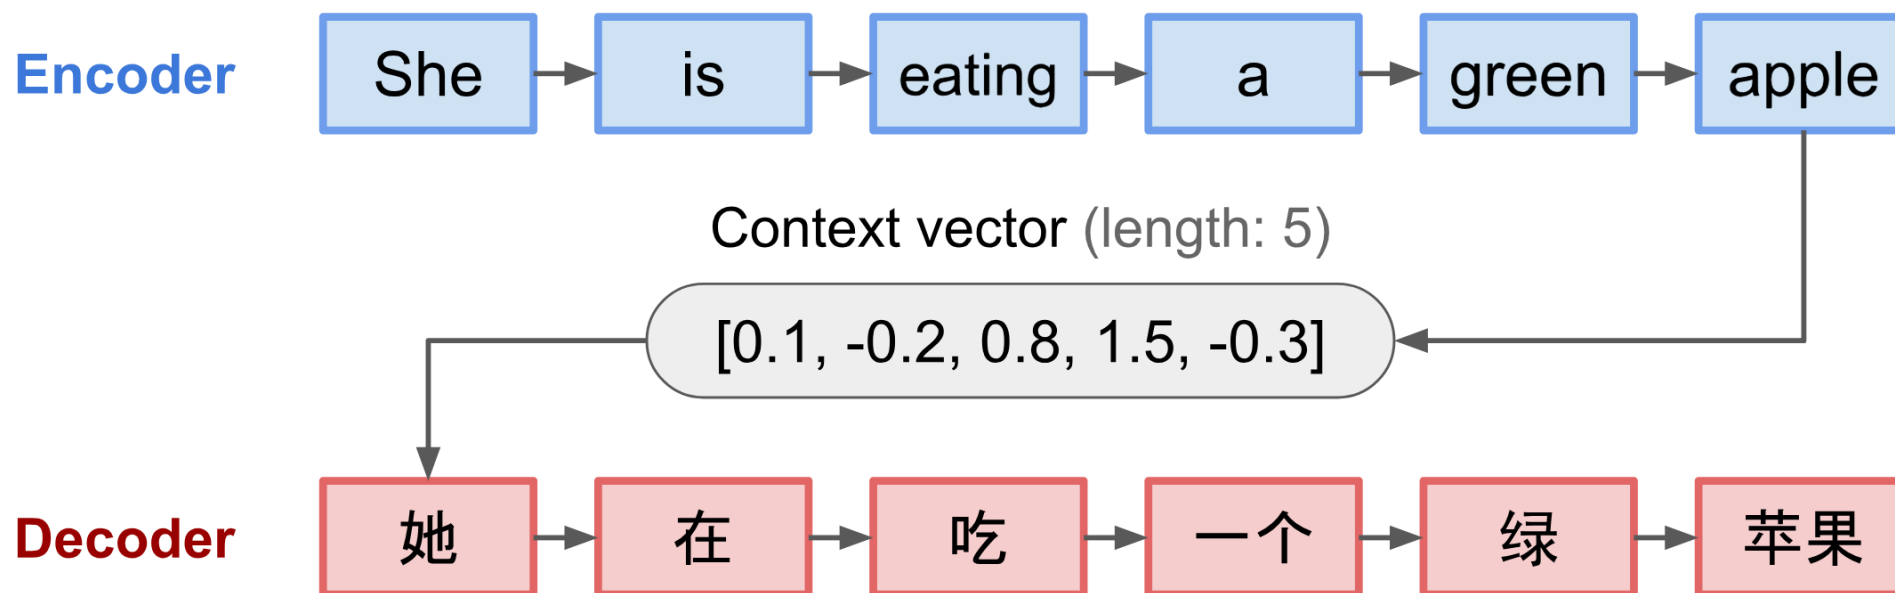
# Seq2Seq Models



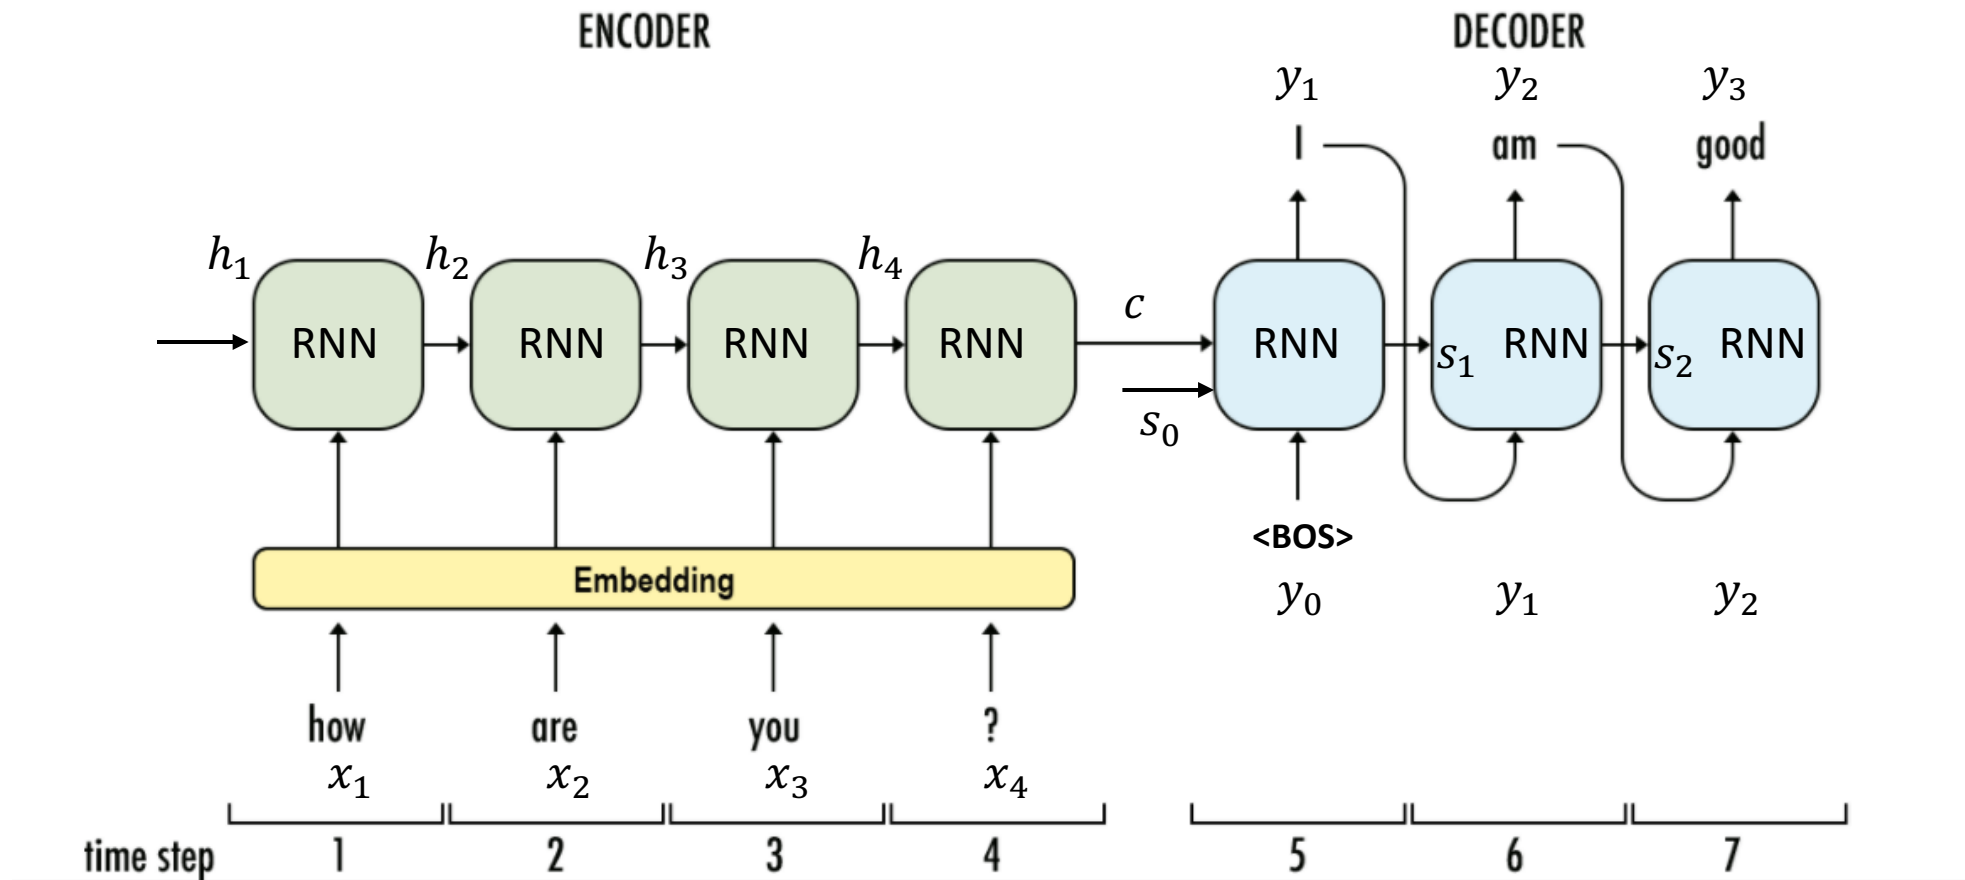
- Note that we require that each sentence ends with a special end-of-sentence symbol **<EOS>**, which enables the model to define a distribution over sequences of all possible lengths.
- We also use a begin-of-sentence symbol **<BOS>** for the decoder.

# Encoder/Decoder


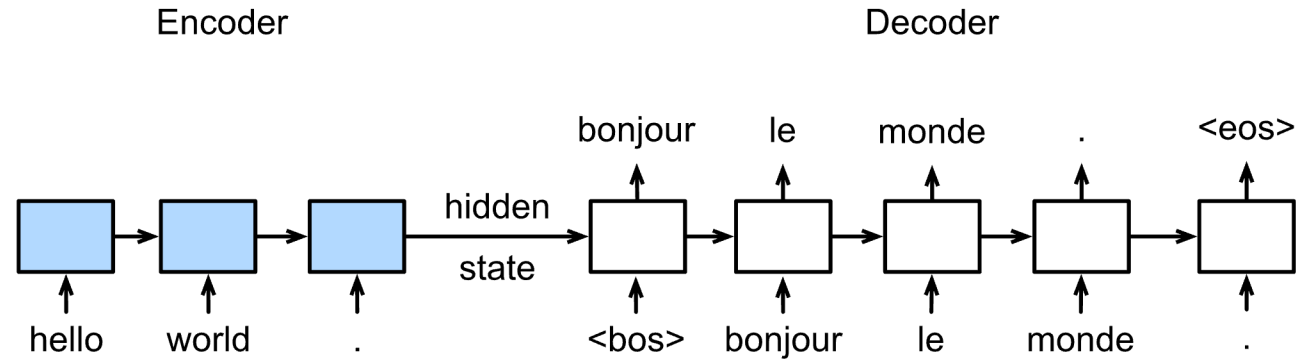
- The encoder-decoder model, translating the sentence “she is eating a green apple” to Chinese.
- The visualization of both encoder and decoder is unrolled in time.



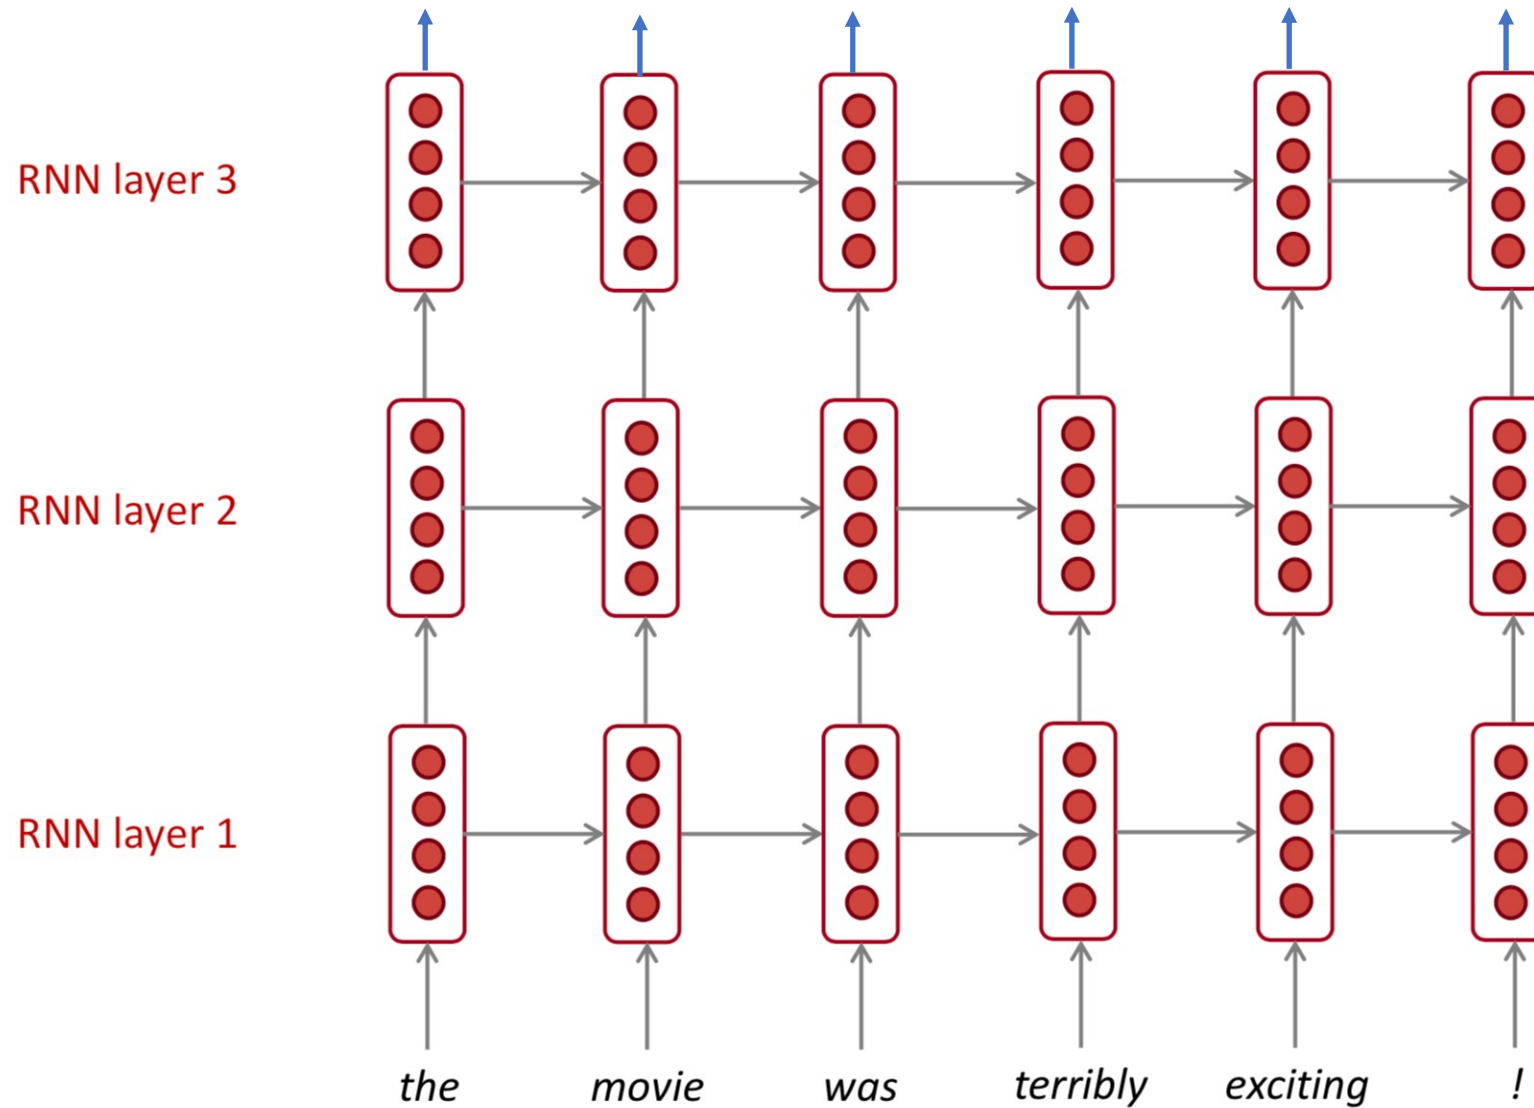
# Encoder/Decoder



# Encoder/Decoder with recurrent layers



# Multi-layer RNNs





# Encoder

- From a probabilistic perspective, the encoder/decoder model is a general method to learn the conditional distribution  $p(y_1, \dots, y_{T'} | x_1, x_2, \dots, x_T)$
- An encoder reads the input sentence, a sequence of vectors  $x = (x_1, x_2, \dots, x_T)$ , into a context vector  $c$
- The most common approach is to use an RNN
  - RNN hidden state:  $h_t = f(x_t, h_{t-1})$
  - Context vector:  $c = q(\{h_1, \dots, h_T\})$

where  $h_t \in \mathbb{R}^n$  is a hidden state at time  $t$ , and  $c$  is a vector generated from the sequence of the hidden states.

- $f$  and  $q$  are some nonlinear functions.

Ilya Sutskever, Oriol Vinyals, Quoc V. Le, 2014, "Sequence to Sequence Learning with Neural Networks, » pp. 3104–311 in NIPS 2014

Kyunghyun Cho, Bart van Merriënboer, Çağlar Gülçehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, Yoshua Bengio:  
Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation. EMNLP 2014: 1724-1734

# Decoder

- The decoder is often trained to predict the next word  $y_{t'}$  given the context vector  $c$  and all the previously predicted words  $\{y_1, y_2, \dots, y_{t'-1}\}$ .
- The decoder defines a probability over the translation  $y = (y_1, \dots, y_{T'})$  by decomposing the joint probability into the ordered conditionals:

$$p(y|x) = \prod_{t'=1}^{T'} p(y_{t'}|y_1, \dots, y_{t'-1}, x) \approx \prod_{t'=1}^{T'} p(y_{t'}|y_1, \dots, y_{t'-1}, \textcolor{red}{c})$$

Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation.

K. Cho, B. van Merriënboer, Ç. Gülçehre, D. Bahdanau, F. Bougares, H. Schwenk, and Y. Bengio. EMNLP, page 1724-1734. ACL (2014)

# Decoder with context vector

- The decoder is also an RNN
  - RNN decoder hidden state:  $s_{t'} = f(y_{t'-1}, s_{t'-1}, c)$
  - The context vector  $c$  can be reused for any hidden state in the decoder

- With an RNN, each conditional probability is modeled as

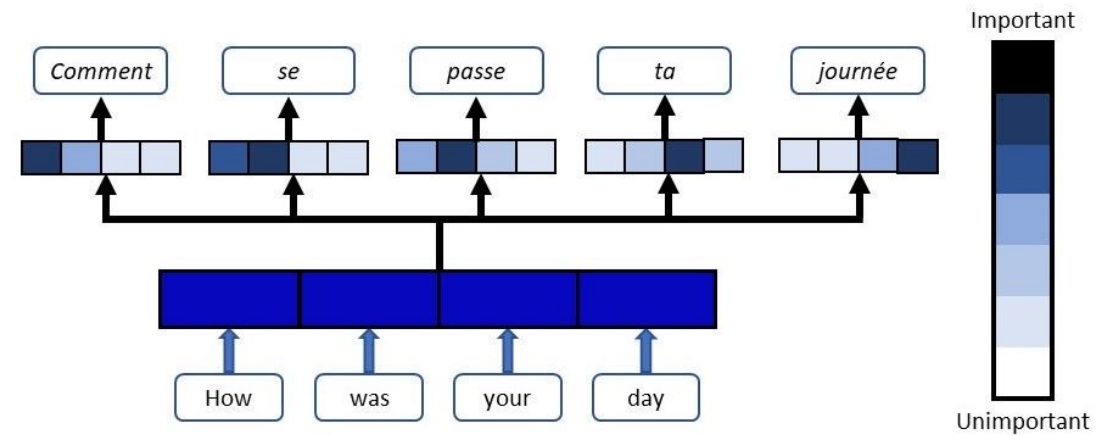
$$p(y_{t'} | y_1, \dots, y_{t'-1}, x) = g(y_{t'-1}, s_{t'}, c)$$

where  $g$  is a nonlinear, potentially multi-layered, function that outputs the probability of  $y_{t'}$ .

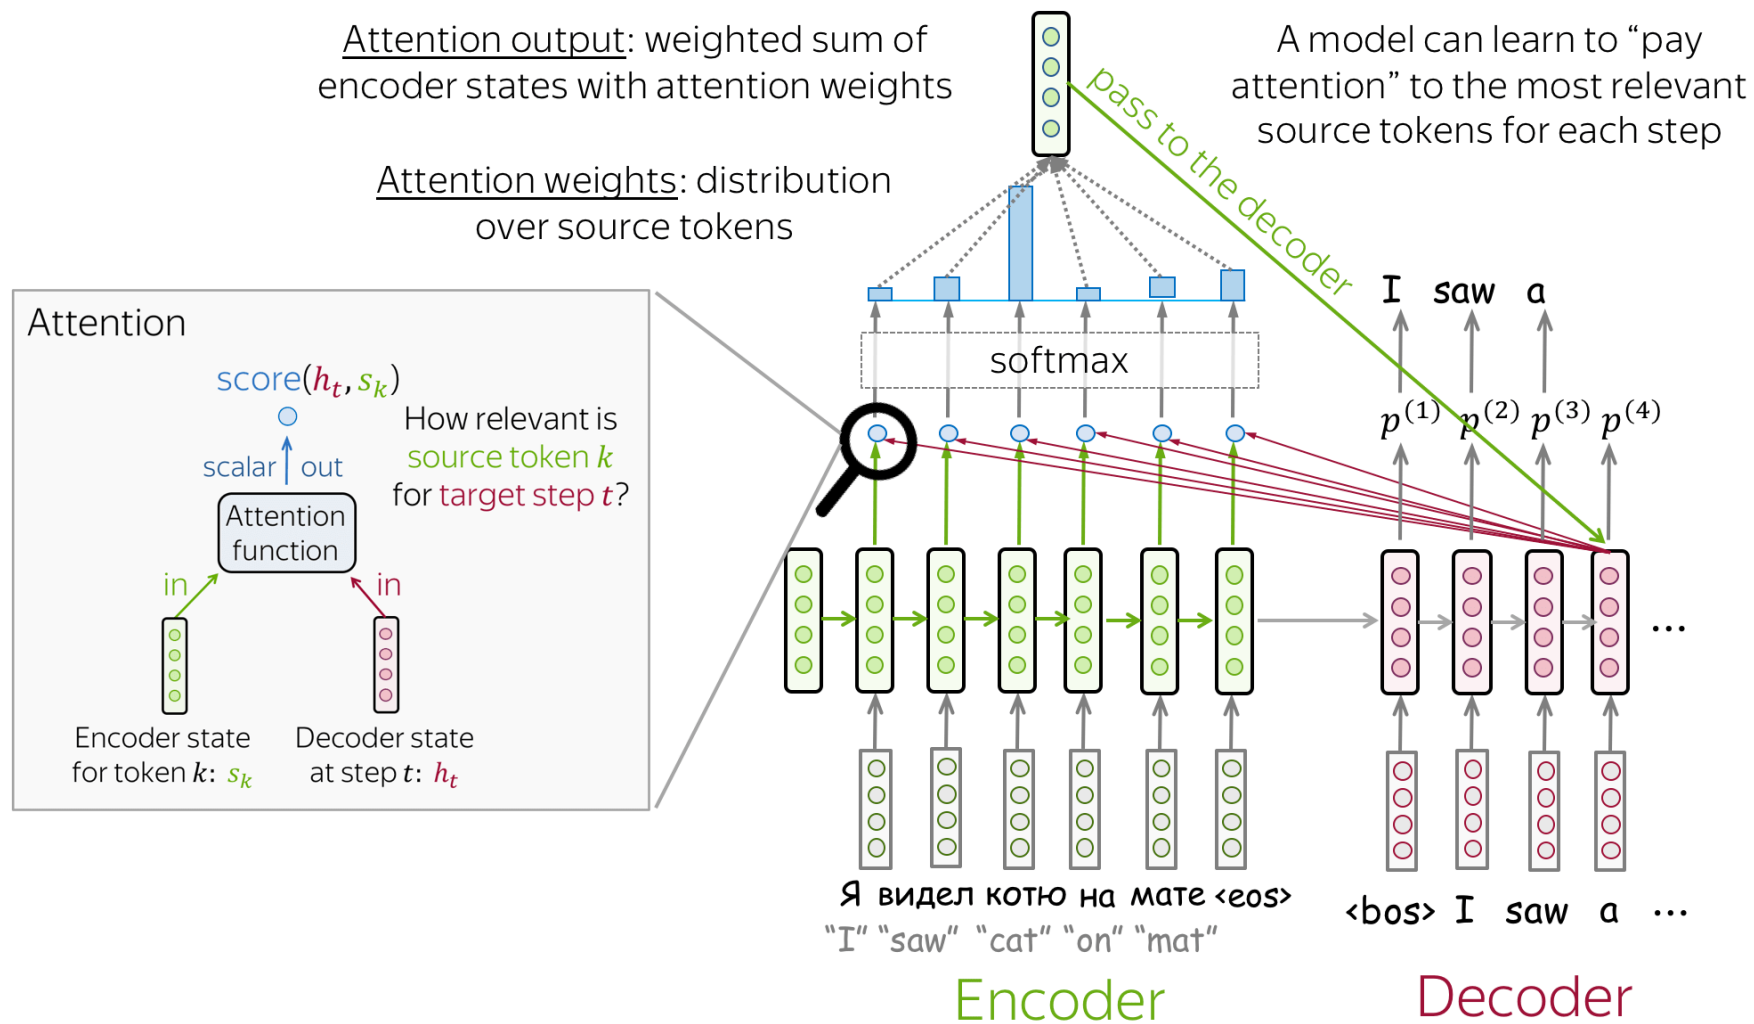
- **Drawback of the single context vector:**
  - A critical and apparent disadvantage of this fixed-length context vector design is incapability of remembering long sentences.
  - Often it has forgotten the first part once it completes processing the whole input.
  - The attention mechanism was born to resolve this problem.

# Attention Mecanism

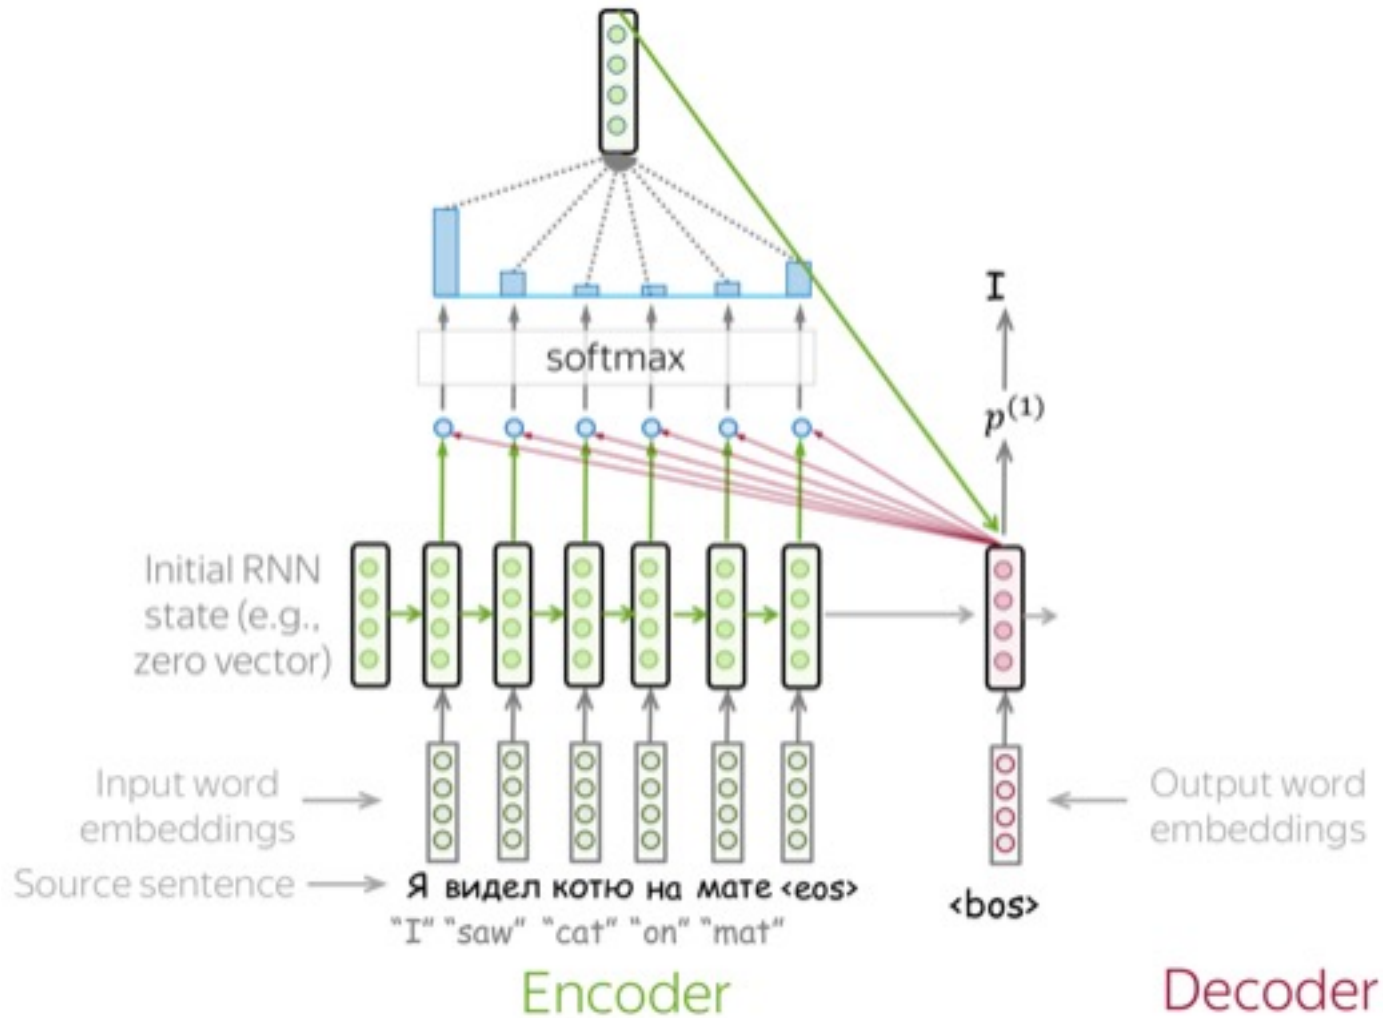
# Decoder with attention




- The decoder is still an RNN
  - RNN decoder hidden state:  $s_{t'} = f(y_{t'-1}, s_{t'-1}, c_{t'})$
- With an RNN, each conditional probability is modeled as
$$p(y_{t'} | y_1, \dots, y_{t'-1}, x) = g(y_{t'-1}, s_{t'}, c_{t'})$$
  - $g$  is a nonlinear, potentially multi-layered, function that outputs the probability of  $y_{t'}$
- Principle: the probability is conditioned on a distinct context vector  $c_{t'}$  for each target word  $y_{t'}$



# An illustration




# An illustration







# An illustration




# An illustration




# An illustration



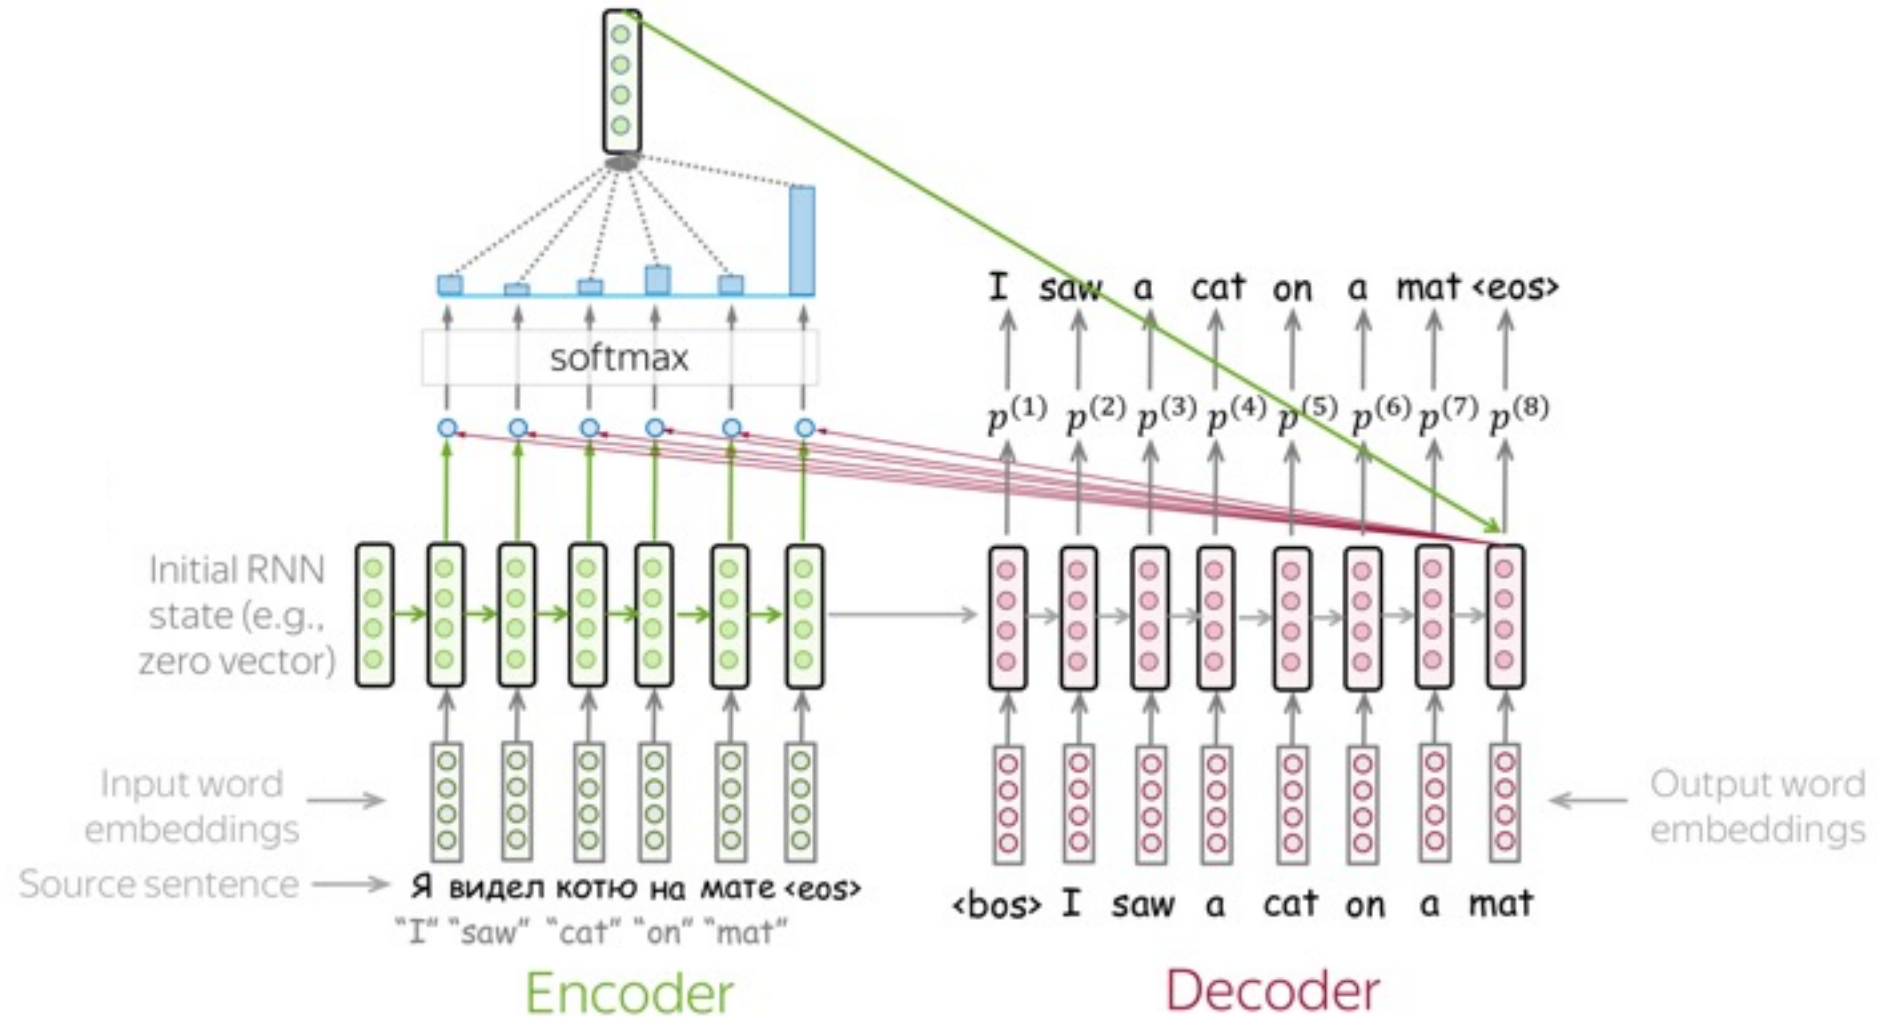
# An illustration



# An illustration



# An illustration



# Annotation

- The context vector  $c_i$  depends on a sequence of annotations  $h = (h_1, h_2, \dots, h_T)$  to which an encoder maps the input sentence.
- Each annotation  $h_i$  contains information about the whole input sequence (especially for a bidirectional RNN) with a strong focus on the parts surrounding the  $i$ -th word of the input sequence.
- Generally, the annotation is a RNN hidden state
- The context vector  $c_i$  is computed as a weighted sum of the annotations  $h_i$ :

$$c_i = \sum_{j=1}^T \alpha_{ij} h_j$$

with  $0 \leq \alpha_{ij} \leq 1$



# Annotation weights

- The weight  $\alpha_{ij}$  of each annotation  $h_j$  is computed by a softmax function

$$\alpha_{ij} = \text{softmax}(e_{ij}) = \frac{\exp(e_{ij})}{\sum_{k=1}^T \exp(e_{ik})}$$

where

$$e_{ij} = \text{score}(s_i, h_j)$$

- The score  $e_{ij}$  measures how well the inputs around position  $j$  and the output at position  $i$  match
- The score is based on the RNN hidden state  $s_i$  (just before emitting  $y_i$ ) and the  $j$ -th annotation  $h_j$  of the input sentence.



# To to compute the score?

- The most popular score are
  - Dot-product:  $e_{ij} = \text{score}(s_i, h_j) = s_i^T h_j = h_j^T s_i$
  - Bilinear function:  $e_{ij} = \text{score}(s_i, h_j) = s_i^T W h_j = h_j^T W s_i$
  - Multi-Layer Perceptron:  $e_{ij} = \text{score}(s_i, h_j) = a(s_i, h_j) = w_2^T \tanh(W_1[s_i, h_j])$
- In the original paper, the authors used the **alignment model**  $a(\cdot)$  as a feedforward neural network which is jointly trained with all the other components of the proposed system.

# Interpretation of $\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^T \exp(e_{ik})}$

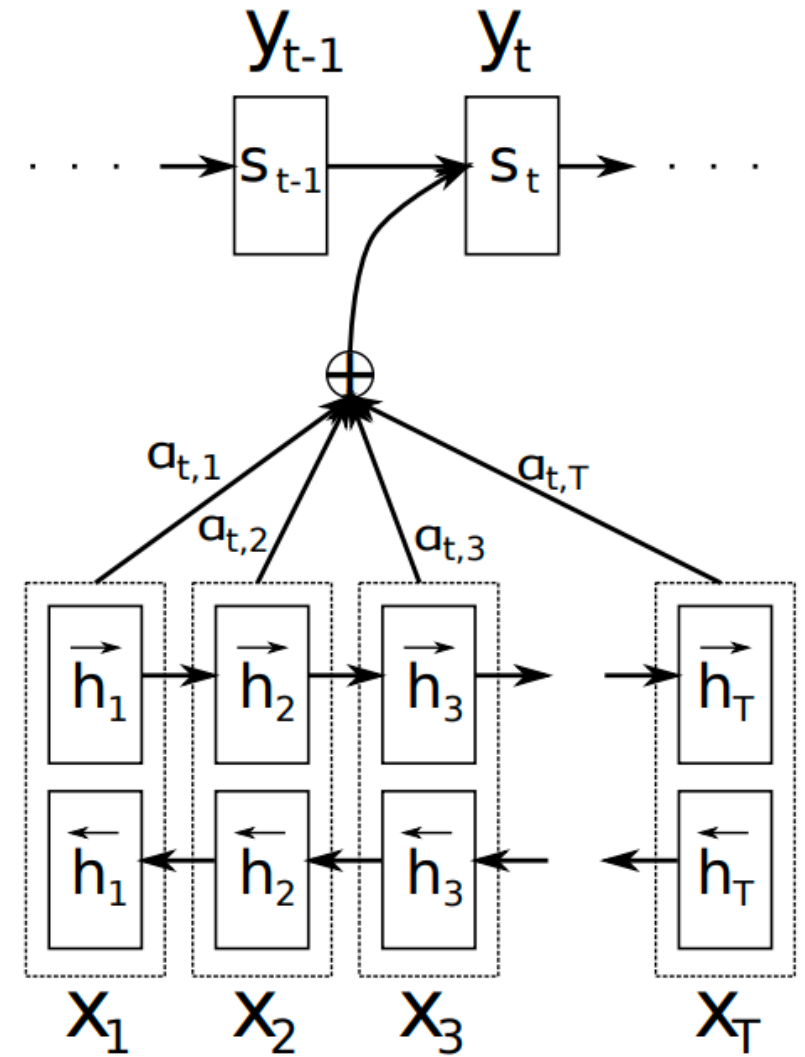
- We can understand the approach of taking a weighted sum of all the annotations as computing an **expected annotation**, where the expectation is over possible alignments.
- Let  $\alpha_{ij}$  be a probability that the target word  $y_i$  is aligned to, or translated from, a source word  $x_j$ . Then, the  $i$ -th context vector  $c_i$  is the expected annotation over all the annotations  $h_j$  with probabilities  $\alpha_{ij}$

$$c_i = \sum_{j=1}^T \alpha_{ij} h_j$$

- The probability  $\alpha_{ij}$ , or its associated score  $e_{ij}$ , reflects the importance of the annotation  $h_j$  with respect to the current decoding state  $s_i$  in deciding the next prediction  $y_i$ .
- Intuitively, this implements a **mechanism of attention** in the decoder.


# Bidirectional RNN

- The hidden state of the encoder is often a bidirectional RNN



# Illustration of the alignment

- The  $x$ -axis and  $y$ -axis of the plot correspond to the words in the source sentence (English) and the generated translation (French), respectively.
- Each pixel shows the weight  $\alpha_{ij}$  of the annotation of the  $j$ -th source word for the  $i$ -th target word, in grayscale (0: black, 1: white).




# Transformer's Encoder

# Transformer's Encoder: Principle


- Self-attention
- Queries, keys and values
- Scaling the dot product
- Multi-head attention

# Example with a sentence



- The word « licked » is most correlated to « cat » (who?) and « owner » (to whom?)

# Comparison with convolution



- The word « licked » is only correlated to words in a given neighborhood (kernel size)



# Self-attention

- Self-attention is a sequence-to-sequence operation:
  - A sequence of vectors goes in, and a sequence of vectors comes out.
  - Let's call the input vectors  $x_1, x_2, \dots, x_t$  and the corresponding output vectors  $y_1, y_2, \dots, y_t$ .
  - The vectors all have dimension  $d$  (the inputs are embedded with an embedding layer).
- To produce output vector  $y_i$ , the self attention operation simply takes a weighted average over all the input vectors

$$y_i = \sum_{j=1}^t w_{ij} x_j$$

where the positive weights  $w_{ij}$  sum to one over all  $j$ .

# Self-attention: basic operation

- The weight

$$w_{ij} = \text{softmax}(e_{ij}) = \text{softmax}(\text{score}(x_i, x_j))$$

is not a parameter, as in a normal neural net, but it is derived from a function over  $x_i$  and  $x_j$ .


- The simplest option for the score function is the dot product:

$$e_{ij} = x_i^T x_j$$

- The dot product gives us a value anywhere between negative and positive infinity, so we apply a softmax to map the values to  $[0,1]$  and to ensure that they sum to 1 over the whole sequence:

$$w_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^t \exp(e_{ik})}$$

# Self-attention



A visual illustration of basic self-attention.  
Note that the softmax operation over the weights is not illustrated.

# Example: How to draw a rose

Query

how to draw a rose

Vidéos :

How To Draw A Rose

YouTube · Art for Kids Hub  
3 févr. 2022

8 temps forts dans cette vidéo

How to Draw a Rose

YouTube · Draw So Cute  
7 févr. 2022

10 temps forts dans cette vidéo

How to Draw + Color a Rose Super EASY Realistic

YouTube · Draw So Cute  
31 janv. 2017

Value

Key

# Queries, keys and values

$$y_i = \sum_{j=1}^t w_{ij} x_j, \quad w_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^t \exp(e_{ik})}, \quad e_{ij} = x_i^T x_j$$

- Every input vector  $x_i$  is used in three different ways in the self attention operation (each role has a name: query, key or value):

$$y_i = \sum_{j=1}^t \frac{\exp(x_i^T x_j)}{\sum_{k=1}^t \exp(x_i^T x_k)} x_j$$

1. **Query**: vector from which the attention is looking for its own output  $y_i$
  2. **Key**: It is compared to every other vector at which the query looks to establish the weights
  3. **Value**: It is used as part of the weighted sum to compute each output vector once the weights have been established.
- In the basic self-attention we've seen so far, each input vector must play all three roles.
  - In the transformer, new vectors for each role are derived, by applying a linear transformation to the original input vector

# Linear transformation for each role

- We can add three  $d \times d$  weight matrices  $W_q$ ,  $W_k$ ,  $W_v$  to compute three linear transformations of each  $x_i$ , for the three different parts of the self attention:

$$q_i = W_q x_i, \quad k_i = W_k x_i, \quad v_i = W_v x_i$$

$$e_{ij} = q_i^T k_j$$

$$w_{ij} = \text{softmax}(e_{ij})$$

$$y_i = \sum_{j=1}^t w_{ij} v_j$$

- This gives the self-attention layer some controllable parameters, and allows it to modify the incoming vectors to suit the three roles they must play.

# Illustration




Illustration of the self-attention with key, query and value transformations.

# Scaling the dot product

- The softmax function can be sensitive to very large input values.
- These kill the gradient, and slow down learning, or cause it to stop altogether.
- Since the average value of the dot product grows with the embedding dimension  $k$ , it helps to scale the dot product back a little to stop the inputs to the softmax function from growing too large:

$$e_{ij} = \frac{q_i^T k_j}{\sqrt{d}}$$


- Why  $\sqrt{d}$ ?
  - Imagine a vector in  $\mathbb{R}^d$  with values all  $c$ :  $(c, c, \dots, c)$ . Its Euclidean length is  $c\sqrt{d}$ .
  - Therefore, we are dividing out the amount by which the increase in dimension increases the length of the average vectors.



# Attention function: matrix form

- In practice, we compute the attention function on a set of queries simultaneously, packed together into a matrix  $Q$  (after the linear transformation if we use it).
  - Initial vectors are the rows of  $Q$
- The keys and values are also packed together into matrices  $K$  and  $V$ .
- We compute the matrix of outputs as:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{Q K^T}{\sqrt{d}}\right) V$$



# Multi-head attention


- Finally, we must account for the fact that a word can mean different things to different neighbours.
  - Consider the following example: « mary gave roses to susan »
  - We see that the word gave has different relations to different parts of the sentence.
    - mary expresses who's doing the giving,
    - roses expresses what's being given,
    - and susan expresses who the recipient is.
- In a single self-attention operation, all this information just gets summed together.
  - If Susan gave Mary the roses instead, the output vector  $y_{gave}$  would be the same, even though the meaning has changed.

# Multi-head attention

- We can give the self attention greater power of discrimination, by combining several self attention mechanisms (which we'll index with  $i$ ), each with different matrices  $W_i^Q$ ,  $W_i^K$ ,  $W_i^V$ . These are called attention heads.
- For input  $x_j$  each attention head produces a different output vector  $y_j^i$ . We concatenate these, and pass them through a linear transformation  $W^O$  to reduce the dimension back to  $d$ .


# Wide self-attention

- If the embedding vector has 256 dimensions, and we have  $h = 8$  attention heads
- For each head, we generate keys, values and queries of 256 dimensions each
- This means that the matrices  $W_i^Q$ ,  $W_i^K$ ,  $W_i^V$  are all  $256 \times 256$  with  $i = 1, \dots, h$



# In practice: narrow self-attention

- The standard option is to cut the embedding vector into  $h$  chunks:
  - If the embedding vector has 256 dimensions, and we have  $h = 8$  attention heads,
  - We cut the embedding vector into 8 chunks of 32 dimensions.
  - For each chunk, we generate keys, values and queries of 32 dimensions each.
  - This means that the matrices  $W_i^Q$ ,  $W_i^K$ ,  $W_i^V$  are all  $32 \times 32$  with  $i = 1, \dots, h$
- The narrow self-attention is faster, and more memory efficient but all else being equal, the wide self-attention does give better results (at the cost of more memory and time).



# Multi-Head Attention (Matrix Form)


- Matrix model:

$$\text{MultiHead}(Q, K, V) = [\text{head}_1; \text{head}_2; \dots, \text{head}_h] W^O$$

$$\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$$


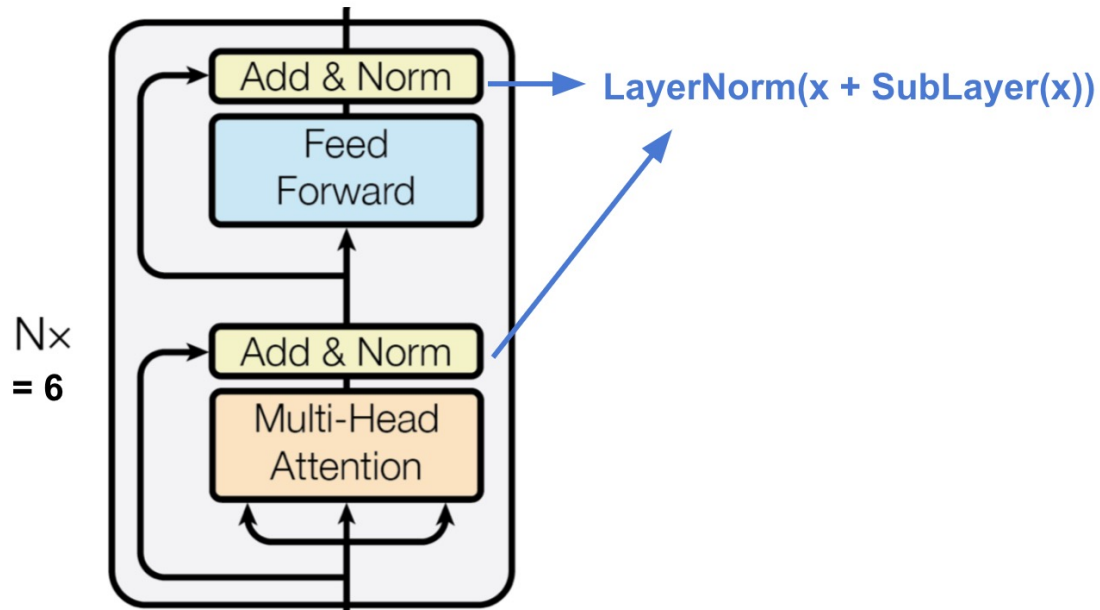
- where

- $W_i^Q \in \mathbb{R}^{d \times d_Q},$
- $W_i^K \in \mathbb{R}^{d \times d_K},$
- $W_i^V \in \mathbb{R}^{d \times d_V},$
- $W^O \in \mathbb{R}^{hd_v \times d}$  are parameter matrices to be learned.



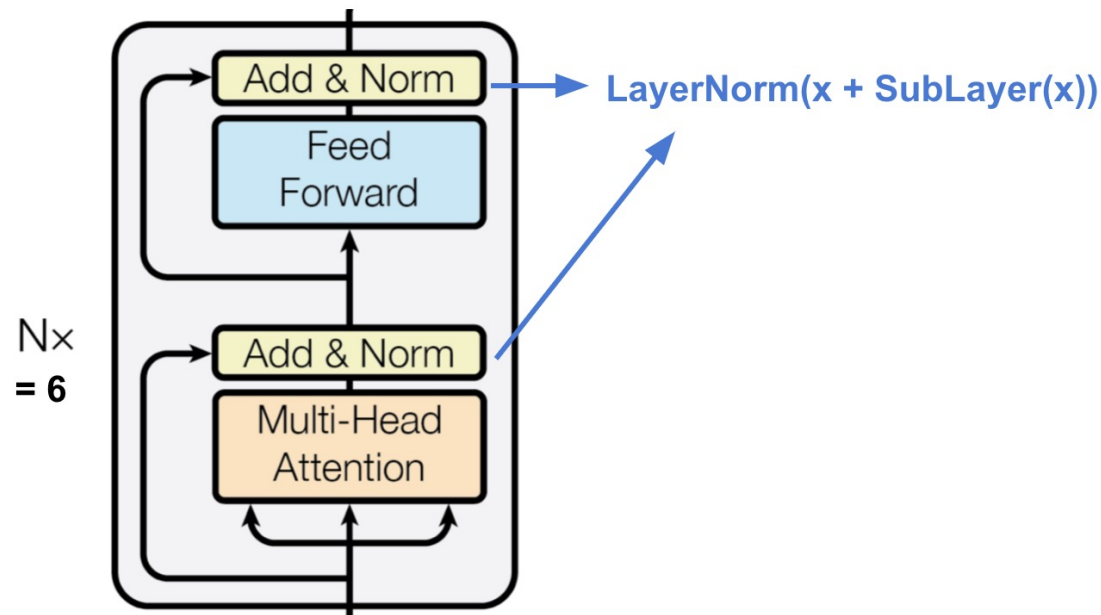
# The Transformer's Encoder

- A stack of  $N = 6$  identical layers.
- Each layer has a multi-head self-attention layer and a simple position-wise fully connected feed-forward network.
  - Attention: "look at other tokens and gather information",
  - FFN: "take a moment to think and process this information").



# The Transformer's Encoder

- Each sub-layer adopts a residual connection and a layer normalization.
  - Equivalent to regularization of the layer
- All the sub-layers output data of the same dimension (for example  $d = 512$ )



Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 770–778, 2016.

Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E Hinton. Layer normalization. *arXiv preprint arXiv:1607.06450*, 2016.



# Block « Add and Norm »

- Layer normalization is similar to batch normalization (but it is not strictly speaking a batch normalization)
- Layer normalization prevents the range of values in the layers from changing too much, which allows faster training and better generalization ability.
- The transformer employs a residual connection around each of the sublayers, followed by layer normalization.
- That is, the output of each sublayer is

$$\text{LayerNorm}(x + \text{Sublayer}(x))$$

where  $\text{Sublayer}(x)$  is the function implemented by the sublayer itself.

# LayerNorm( $x + \text{Sublayer}(x)$ )

- For a batch  $\{x_n\}_{n=1,\dots,N}$  of  $N$  vectors  $x_n \in \mathbb{R}^K$ , also written as  $\{x_{n,k}\} \in \mathbb{R}^{N \times K}$ , the expectation and variance accros spatial dimensions are « estimated » by

$$\mu_n = \frac{1}{K} \sum_{k=1}^K x_{n,k} \in \mathbb{R}, \quad \sigma_n^2 = \frac{1}{K} \sum_{k=1}^K (x_{n,k} - \mu_n)^2 \in \mathbb{R}$$

- Layer Normalization (LayerNorm in Pytorch)

$$\hat{x}_{n,k} = \frac{x_{n,k} - \mu_n}{\sqrt{\sigma_n^2 + \epsilon}} \in \mathbb{R} \quad \Rightarrow \quad \hat{x}_n = \begin{pmatrix} \hat{x}_{n,1} \\ \vdots \\ \hat{x}_{n,K} \end{pmatrix} \in \mathbb{R}^K$$

$$LN_{\gamma,\beta}(x_n) = \gamma \hat{x}_n + \beta$$

- $\gamma$  and  $\beta$  are learnable affine transform parameters

# Input encoding

# Encoding the inputs

- The initial inputs of the encoder are the embeddings of the sequence of inputs (typically Word Embeddings)
- The initial inputs of the decoder are the embeddings of the outputs.
- The order of sequence (position of words in a sentence) information is very important.
  - Since there is no recurrence, this information on the absolute (or relative) position in a sequence is represented by the use of “position encodings”.

# Positional Encodings

- Recurrent Neural Networks (RNNs) inherently take the order of word into account
  - They parse a sentence word by word in a sequential manner.
- But the Transformer architecture ditched the recurrence mechanism in favor of multi-head self-attention mechanism. Avoiding the RNNs' method of recurrence will result in massive speed-up in the training time.
- As each word in a sentence simultaneously flows through the Transformer's encoder/decoder stack, the model itself doesn't have any sense of position/order for each word.
- One possible solution to give the model some sense of order is to add a piece of information to each word about its position in the sentence. We call this « piece of information », the positional encoding.

Full input of a transformer = Embedding vector + Positional encoding

# Criteria for Positional Encodings

- The first idea that might come to mind is to assign a number to each time-step within the  $[0, 1]$  range in which 0 means the first word and 1 is the last time-step.
  - One of the problems it will introduce is that you can't figure out how many words are present within a specific range. In other words, time-step doesn't have consistent meaning across different sentences.
- Another idea is to assign a number to each time-step linearly: the first word is given "1", the second word is given "2", and so on.
  - The problem with this approach is that not only the values could get quite large, but also our model can face sentences longer than the ones in training.
  - In addition, our model may not see any sample with one specific length which would hurt generalization of our model.
- Ideally, the following criteria should be satisfied:
  - It should output a unique encoding for each time-step (word's position in a sentence)
  - Distance between any two time-steps should be consistent across sentences with different lengths.
  - The model should generalize to longer sentences without any efforts. Its values should be bounded.
  - It must be deterministic.

# Proposed Method for Transformer


- The encoding proposed by the authors satisfies all of those criteria.
  - First of all, it isn't a single number. Instead, it's a  $d$ -dimensional (same dimension as word embedding) vector  $p_t$  that contains information about a specific position  $t$  in a sentence.
  - Secondly, this encoding is not integrated into the model itself. Instead, this vector is used to equip each word with information about its position in a sentence.
  - According to the authors, for any fixed offset  $s$ ,  $p_{t+s}$  can be represented as a linear function of  $p_t$
- Let  $t$  the desired position in an input sentence,  $p_t = (p_t(0), \dots, p_t(d-1)) \in \mathbb{R}^d$  be its corresponding encoding. Then,

$$p_t(i) = \begin{cases} \sin(w_k t) & \text{if } i = 2k \\ \cos(w_k t) & \text{if } i = 2k + 1 \end{cases} \quad \text{with} \quad w_k = \frac{1}{10000^{2k/d}}$$

- Example:  $p_t(0) = \sin\left(\frac{t}{10000^{0/d}}\right)$ ,  $p_t(1) = \cos\left(\frac{t}{10000^{0/d}}\right)$ ,  $p_t(2) = \sin\left(\frac{t}{10000^{2/d}}\right)$ ,  $p_t(3) = \cos\left(\frac{t}{10000^{2/d}}\right)$ , ...

# Visualizing the Positional Encodings

- The 128-dimensional positional encoding for a sentence with the maximum length of 50.
- Each row represents the embedding vector  $p_t$





# Interpretation as Circles

- You could see this encoding as  $d/2$  pairs of circles rather than  $d$  dimensional vectors.
- When you fix the depth  $i = 2k$  of each encoding, you can extract a 2 dimensional vector

$$PE_i(t) = \left( \sin\left(\frac{t}{10000^{i/d}}\right), \cos\left(\frac{t}{10000^{i+1/d}}\right) \right)$$

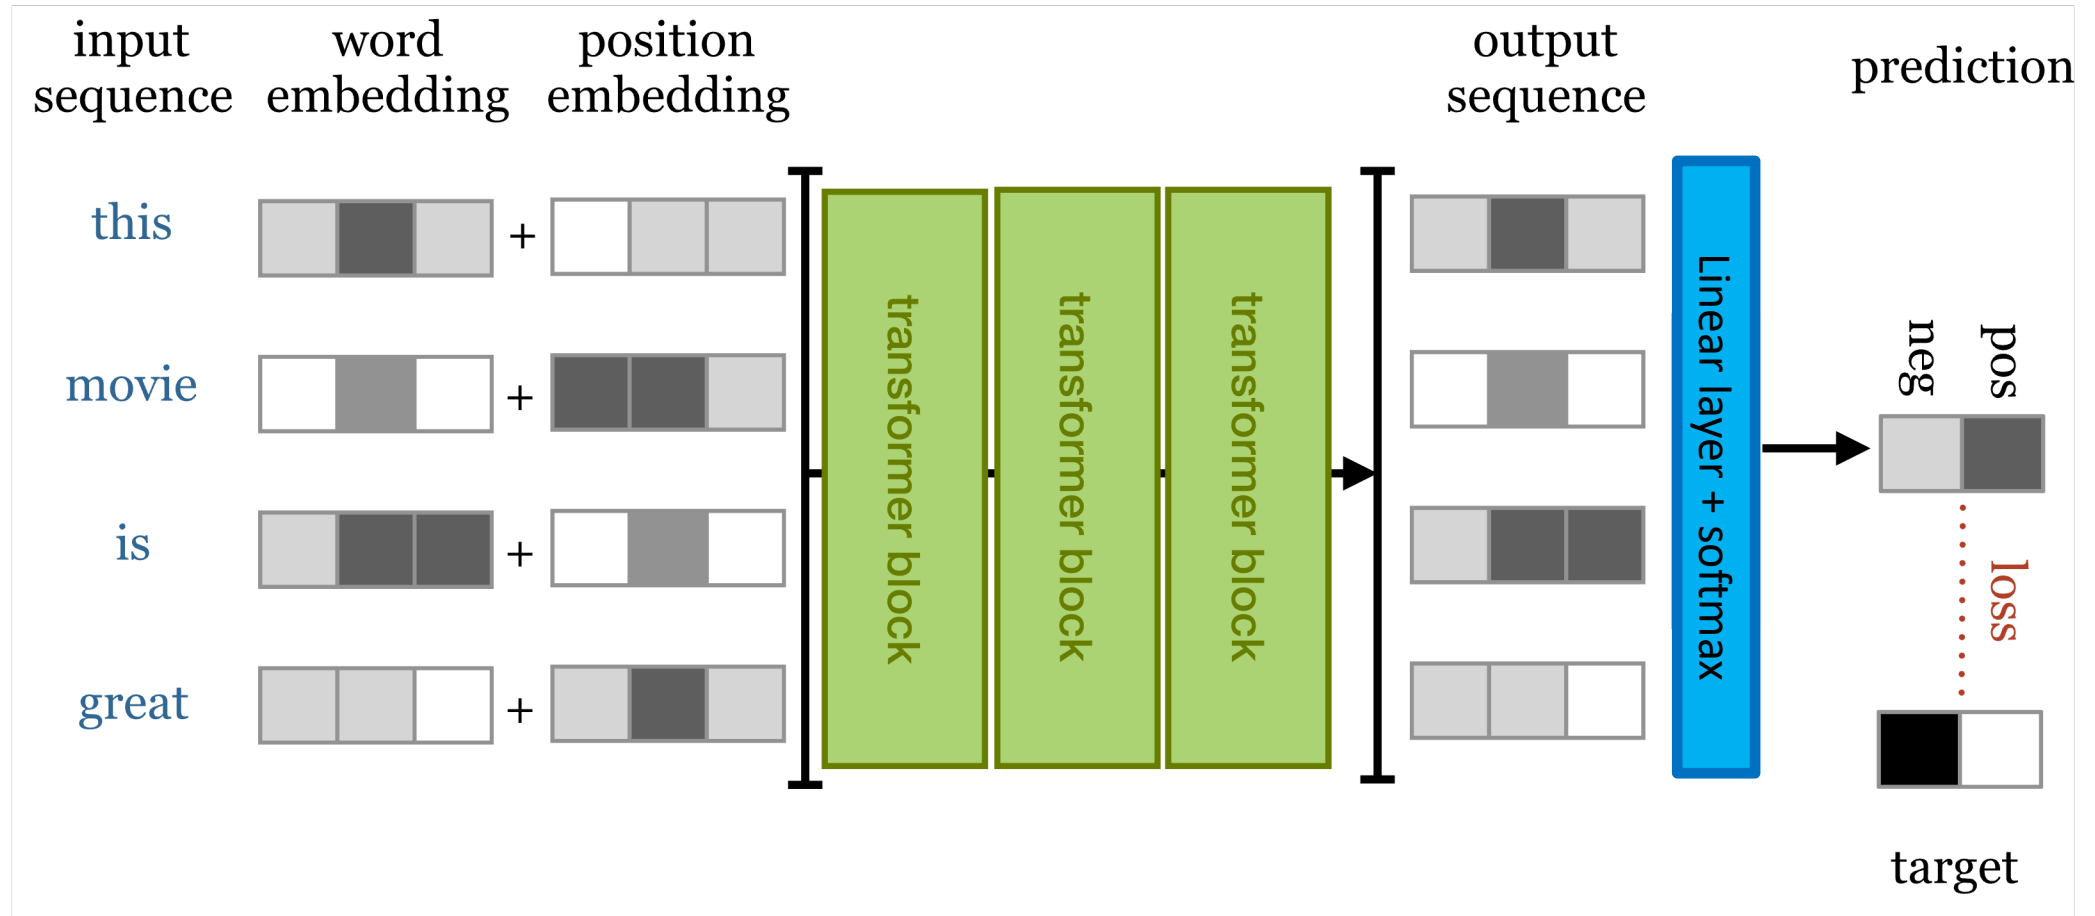
- If you constantly change the position value  $t$ , the vector  $PE_i(t)$  rotates clockwise on the unit circle.
- The deeper the dimension  $i$  of the embedding is, the smaller the frequency of rotation is.
- Note that a linear transformation (a rotation) allows us to change  $PE_i(t)$  into  $PE_i(t + s)$  for any offset  $s$

# Classification transformer


# Classification transformer

- The simplest transformer we can build is a sequence classifier.
- The heart of the architecture will simply be a large chain of transformer blocks.
- All we need to do is work out how to feed it the input sequences, and how to transform the final output sequence into a single classification.

# Producing a classification



# Transformer block



# Transformer block

- The block applies, in sequence:
  - a self attention layer,
  - a layer normalization,
  - a feed forward layer (a single MLP applied independently to each vector),
  - and another layer normalization.
  - residual connections are added around both, before the normalization.
- The order of the various components is not set in stone

# Transformer block in Pytorch

```
class TransformerBlock(nn.Module):  
    def __init__(self, k, heads):  
        super().__init__()  
        self.attention = SelfAttention(k, heads=heads)  
        self.norm1 = nn.LayerNorm(k)  
        self.norm2 = nn.LayerNorm(k)  
        self.ff = nn.Sequential(  
            nn.Linear(k, 4 * k),  
            nn.ReLU(),  
            nn.Linear(4 * k, k)  
        )  
  
    def forward(self, x):  
        attended = self.attention(x)  
        x = self.norm1(attended + x)  
        feedforward = self.ff(x)  
        return self.norm2(feedforward + x)
```

k: input size (word embedding + positional encoding)

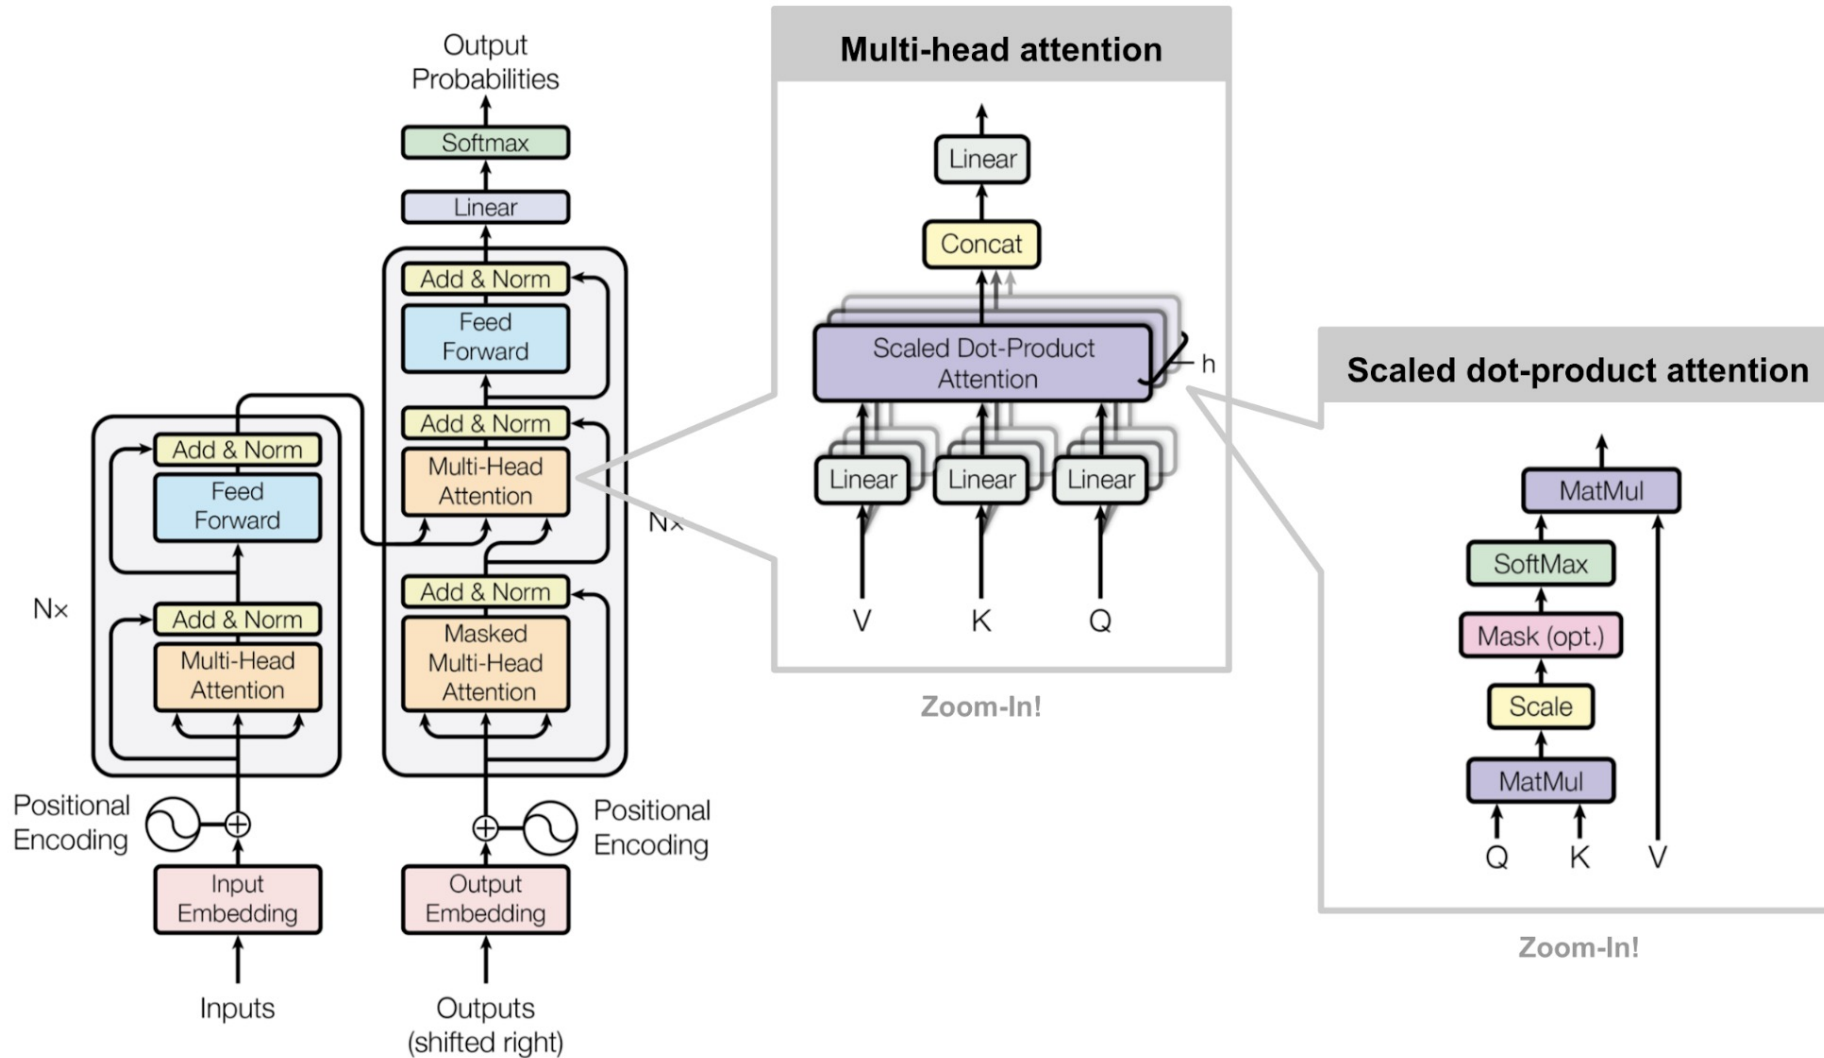
# Transformer for sequence transduction



# Transformer for sequence transduction

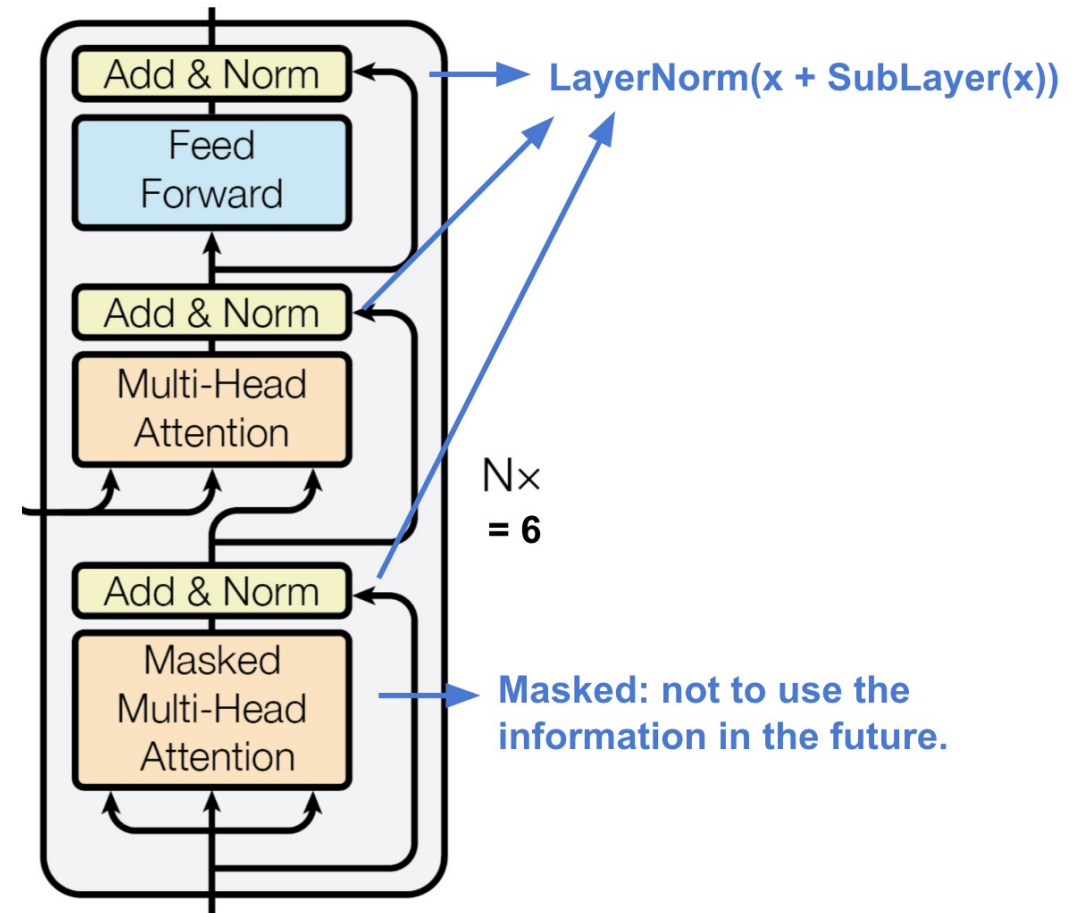
- The encoder maps an input sequence of symbol representations  $x = (x_1, x_2, \dots, x_T)$  to a sequence of continuous representations  $z = (z_1, z_2, \dots, z_T)$ .
- Given  $z$ , the decoder then generates an output sequence  $y = (y_1, y_2, \dots, y_{T'})$  of symbols one element at a time.
- At each step the model is auto-regressive, consuming the previously generated symbols as additional input when generating the next.

# Full Architecture




# The Transformer's Decoder

- A stack of  $N = 6$  identical layers
- Each layer has two sub-layers of multi-head attention mechanisms and one sub-layer of fully-connected feed-forward network.
- Similar to the encoder, each sub-layer adopts a residual connection and a layer normalization.
- The first multi-head attention sub-layer is modified to prevent positions from attending to subsequent positions, as we don't want to look into the future of the target sequence when predicting the current position.



# Masked Multi-head Attention

- With a transformer, the output depends on the entire input sequence, so prediction of the next character becomes vacuously easy, just retrieve it from the input.
- To use self-attention as an autoregressive model, we'll need to ensure that it cannot look forward into the sequence. We do this by applying a mask to the matrix of dot products, before the softmax is applied.
- This mask disables all elements above the diagonal of the matrix.



- Note that the multiplication symbol is slightly misleading: we actually set the masked out elements (the white squares) to  $-\infty$  to get very negative raw weights, then the softmax returns 0

# Conclusion

# Conclusion

- Dealing with sequence of unknown size is a hot topic
- Attention is nowadays a crucial mechanism
- Transformers are exploiting self-attention
- Transformers are powerful and generic
- A transformer is a significant alternative to a recurrent neural network
- Currently extended to computer vision

# Train XLNet (2020)

- XLNet is a language model developed by CMU and Google Research which outperforms the previous SOTA model BERT (Bidirectional Encoder Representations from Transformers) on 20 language tasks including SQuAD, GLUE, and RACE; and has achieved SOTA (State Of The Art) results on 18 of these tasks.
- The authors said: “We train XLNet-Large on 512 TPU v3 chips for 500K steps with an Adam optimizer, linear learning rate decay and a batch size of 2048, which takes about 2.5 days.”
- A Cloud TPU v3 device, which costs US\$8 per hour on Google Cloud Platform, has eight independent embedded chips.
- The total training cost should be  $512 \text{ (chips)} * (\text{US\$8}/8) \text{ (cost per chip)} * 24 \text{ (hours)} * 2.5 \text{ (days)} = \$30,720.$
- Hence, the training cost might be significant for SOTA methods.